MIPS
Assembly Language Programming
using QtSpim

Ed Jorgensen, Ph.D.

Version 1.1.50

July 2019

Cover image:
    MIPS R3000 Custom Chip
    http://commons.wikimedia.org/wiki/File:RCP-NUS_01.jpg

Table of Contents

Table of Contents

Table of Contents

1.0   Introduction

There are a number of excellent, comprehensive, and in-depth texts on MIPS assembly language programming. This is not one of them.

The purpose of this text is to provide a simple and free reference for university level programming and architecture units that include a brief section covering MIPS assembly language programming. The text assumes usage of the QtSpim simulator. An appendix is included that covers the download, installation, and basic use of the QtSpim simulator.

The scope of this text addresses basic MIPS assembly language programming including instruction set usage, stacks, procedure/function calls, QtSpim simulator system services, multiple dimension arrays, and basic recursion.

1.1 Additional References

Some key references for additional information are listed below:

- *MIPS Assembly-language Programmer Guide,* Silicon Graphics
- *MIPS Software Users Manual,* MIPS Technologies, Inc.
- *Computer Organization and Design: The Hardware/Software Interface,* Hennessy and Patterson

More information regarding these references can be found on the Internet.

Chapter 1.0 ◄ Introduction

2.0 MIPS Architecture Overview

This chapter presents a basic, general overview of the architecture of the MIPS processor.

The MIPS architecture is a Reduced Instruction Set Computer (RISC). This means that there is a smaller number of instructions that use a uniform instruction encoding format. Each instruction/operation does one thing (memory access, computation, conditional, etc.). The idea is to make the lesser number of instructions execute faster. In general RISC architectures, and specifically the MIPS architecture, are designed for high-speed implementations.

2.1 Architecture Overview

The basic components of a computer include a Central Processing Unit (CPU), Primary Storage or Random Access Memory (RAM), Secondary Storage (i.e., Disk Drive, SSD, etc.), Input/Output devices (i.e., screen and keyboard), and an interconnection referred to as BUS. A very basic diagram of a computer architecture is as follows:

Illustration 1: Computer Architecture

Programs and data are typically stored on secondary storage (i.e. SSD, disk drive).

When a program is executed, it must be copied from the disk drive into the RAM memory. The CPU executes the program from RAM. This is similar to storing a term paper on the disk drive, and when writing/editing the term paper, it is copied from the disk drive into memory. When done, the updated version is stored back to the disk drive.

2.2 Data Types/Sizes

The basic data types include integer, floating-point, and characters.

This architecture supports data storage sizes of byte, halfword (sometimes referred to as just *half*), or word sizes. Floating-point must be of either word (32-bit) size or double word (64-bit) size. Character data is typically a byte and a string is a series of sequential bytes.

The MIPS architecture supports the following data/memory sizes:

Name	Size
byte	8-bit integer
halfword	16-bit integer
word	32-bit integer
float	32-bit floating-point number
double	64-bit floating-point number

The halfword is often referred to as just '*half*'. Lists or arrays (sets of memory) can be reserved in any of these types. In addition, an arbitrary number of bytes can be defined with the ".space" directive.

2.3 Memory

Memory can be viewed as a series of bytes, one after another. That is, memory is *byte addressable*. This means each memory address holds one byte of information. To store a word, four bytes are required which use four memory addresses.

Additionally, the MIPS architecture as simulated in QtSpim is **little-endian**. This means that the Least Significant Byte (LSB) is stored in the lowest memory address. The Most Significant Byte (MSB) is stored in the highest memory location.

For a word (32-bits), the MSB and LSB are allocated as shown below.

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
MSB																											**LSB**				

For example, assuming the following declarations:

```
num1:    .word    42
num2:    .word    5000000
```

Recall that 42_{10} in hex, word size, is 0x0000002A and $5,000,000_{10}$ in hex, word size, is 0x004C4B40.

For a little-endian architecture, the memory picture would be as follows:

variable name	value	address
	?	0x100100C
	00	0x100100B
	4C	0x100100A
	4B	0x1001009
Num2 →	40	0x1001008
	00	0x1001007
	00	0x1001006
	00	0x1001005
Num1 →	2A	0x1001004
	?	0x1001003

Based on the little-endian architecture, the LSB is stored in the lowest memory address and the MSB is stored in the highest memory location.

2.4 Memory Layout

The general memory layout for a program is as shown:

high memory	**stack**
	heap
	uninitialized data
	data
	text (code)
low memory	**reserved**

The reserved section is not available to user programs. The text (or code) section is where the machine language (i.e., the 1's and 0's that represent the code) is stored. The data section is where the initialized data is stored. This includes declared variables that have been provided an initial value at assemble time. The uninitialized data section is where declared variables that have not been provided an initial value are stored. If accessed before being set, the value will not be meaningful. The heap is where dynamically allocated data will be stored (if requested). The stack starts in high memory and grows downward.

The QtSpim simulator does not distinguish between the initialized and uninitialized data sections. Later sections will provide additional detail for the text and data sections.

2.5 CPU Registers

A CPU register, or just register, is a temporary storage or working location built into the CPU itself (separate from memory). Computations are typically performed by the CPU using registers.

The MIPS has 32, 32-bit integer registers (**$0** through **$31**) and 32, 32-bit floating-point registers (**$f0** through **$f31**). Some of the integer registers are used for special purposes. For example, **$29** is dedicated for use as the stack pointer register, referred to as **$sp**.

The registers available and typical register usage is described in the following table.

Register Name	Register Number	Register Usage
$zero	$0	Hardware set to 0
$at	$1	Assembler temporary
$v0 - $v1	$2 - $3	Function result (low/high)
$a0 - $a3	$4 - $7	Argument Register 1
$t0 - $t7	$8 - $15	Temporary registers
$s0 - $s7	$16 - $23	Saved registers
$t8 - $t9	$24 - $25	Temporary registers
$k0 - $k1	$26 - $27	Reserved for OS kernel
$gp	$28	Global pointer
$sp	$29	Stack pointer
$fp	$30	Frame pointer
$ra	$31	Return address

The register names convey specific usage information. The register names will be used in the remainder of this document. Further sections will expand on register usage conventions and address the 'temporary' and 'saved' registers.

2.5.1 Reserved Registers

The following reserved registers should not be used in user programs.

Register Name	Register Usage
$k0 - $k1	Reserved for use by the Operating System
$at	Assembler temporary
$gp	Global pointer
$epc	Exception program counter

The **$k0** and **$k1** registers are reserved for use by the operating system and should not be used in user programs. The **$at** register is used by the assembler and should not be used in user programs. The **$gp** register is used as a pointer to global data (as needed) and should not be used in user programs.

2.5.2 Miscellaneous Registers

In addition to the previously listed registers, there are some miscellaneous registers which are listed in the table:

Register Name	Register Usage
$pc	Program counter
$status or $psw	Status Register
$cause	Exception cause register
$hi	Used for some multiple/divide operations
$lo	

The **$pc** or program counter register points to the next instruction to be executed and is automatically updated by the CPU after each instruction is executed. This register is not typically accessed directly by user programs.

The **$status** or status register, also called **$psw**, is the processor status register and is updated after each instruction by the CPU. This register is not typically directly accessed by user programs.

The **$cause** or exception cause register is used by the CPU in the event of an exception or unexpected interruption in program control flow. Examples of exceptions include division by 0, attempting to access an illegal memory address, or attempting to execute an invalid instruction (e.g., trying to execute a data item instead of code).

The **$hi** and **$lo** registers are used by some specialized multiply and divide instructions. For example, a multiple of two 32-bit values can generate a 64-bit result, which is stored in **$hi** and **$lo** (32-bits each or a total of 64-bits).

2.6 CPU / FPU Core Configuration

The following diagram shows a basic configuration of the MIPS processor internal architecture.

MIPS Chip Core Configuration

CPU	FPU
integer operations $0 - $31	float operations $f0 - $f31
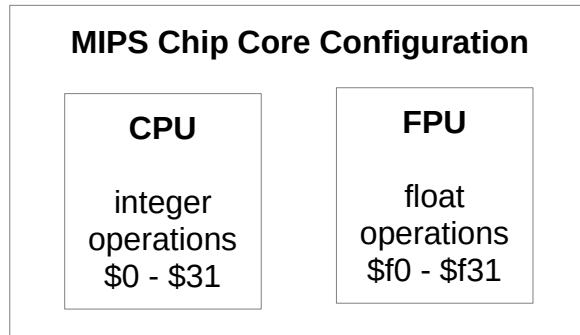
The FPU (floating-point unit) is also referred to as the FPU co-processor or simply co-processor 1.

Chapter 2.0 ◄ MIPS Architecture Overview

3.0 Data Representation

Data representation refers to how information is stored within the computer. There is a specific method for storing integers which is different than storing floating-point values which is different than storing characters. This chapter presents a brief summary of the integer, floating-point, and ASCII representation schemes. It is assumed the reader is already generally familiar with the binary, decimal, and hexadecimal numbering systems.

3.1 Integer Representation

Representing integer numbers refers to how the computer stores or represents a number in memory. As you know, the computer represents numbers in binary. However, the computer has a limited amount of space that can be used for each number or variable. This directly impacts the size, or range, of the number that can be represented. For example, a byte (8 bits) can be used to represent 2^8 or 256 different numbers. Those 256 different numbers can be *unsigned* (all positive) in which case we can represent any number between 0 and 255 (inclusive). If we choose *signed* (positive and negative), then we can represent any number between -128 and +127 (inclusive).

If that range is not large enough to handle the intended values, a larger size must be used. For example, a halfword (16 bits) can be used to represent 2^{16} or 65,536 different numbers, and a word can be used to represent 2^{32} or 4,294,967,296 different numbers. So, if you wanted to store a value of 100,000 then a word would be required.

The following table shows the ranges associated with typical sizes:

Size	Size	Unsigned Range	Signed Range
Bytes (8 bits)	2^8	0 to 255	-128 to +127
Halfwords (16 bits)	2^{16}	0 to 65,535	$-32{,}768$ to $+32{,}767$
Words (32 bits)	2^{32}	0 to 4,294,967,295	$-2{,}147{,}483{,}648$ to $+2{,}147{,}483{,}647$

Chapter 3.0 ◄ Data Representation

In order to determine if a value can be represented, you will need to know the size of storage element (byte, halfword, word) being used and if the values are signed or unsigned values.

- For representing *unsigned* values within the range of a given storage size, standard binary is used.
- For representing *signed* values within the range, **two's complement** is used. Specifically, the two's complement encoding process applies to the values in the negative range. For values within the positive range, standard binary is used.

Additional detail regarding two's complement is provided in the next section.

For example, the unsigned byte range can be represented using a number line as follows:

|—————————————————|—————————————————————————|

0   **255**

For example, the signed byte range can also be represented using a number line as follows:

|————————————————|————————————————|

-128   **0**   **+127**

The same concept applies to halfwords and words with larger ranges.

Unsigned values have a different, positive only, range. The range of the signed value has some overlap with the unsigned values. For example, when the unsigned and signed values are within the overlapping positive range (0 to +127):

- A signed byte representation of 12 is $0x0C_{16}$
- An unsigned byte representation of -12 is also $0x0C_{16}$

When the unsigned and signed values are outside the overlapping range:

- A signed byte representation of -15 is $0xF1_{16}$
- An unsigned byte representation of 241 is also $0xF1_{16}$

This overlap can cause confusion unless the data types are clearly and correctly defined.

3.1.1 Two's Complement

The following describes how to find the two's complement representation for negative values.

To take the two's complement of a number:

1. Take the one's complement (negate)
2. Add 1 (in binary)

The same process is used to encode a decimal value into two's complement and from two's complement back to decimal. The following sections provide some examples.

3.1.2 Byte Example

For example, to find the byte size, two's complement representation of -9 and -12.

9 (8+1) =	00001001		12 (8+4) =	00001100
Step 1	11110110		Step 1:	11110011
Step 2	11110111			11110100
-9 (in hex) =	F7		-12 (in hex) =	F4

Note, all bits for the given size, byte in this example, must be specified.

3.1.3 Halfword Example

To find the halfword size, two's complement representation of -18 and -40.

18 (16+2) =	0000000000010010		40 (32+8) =	0000000000101000
Step 1	1111111111101101		Step 1	1111111111010111
Step 2	1111111111101110		Step 2	1111111111011000
-18 (hex) =	FFEE		-40 (hex) =	FFD8

Note, all bits for the given size, halfwords in these examples, must be specified.

3.2 Unsigned and Signed Addition

As previously noted, the unsigned and signed representations may provide different interpretations for the final value being represented. However, the addition and subtraction operations are the same. For example:

241	11110001		-15	11110001
+ 7	00000111		+ 7	00000111
248	11111000		-8	11111000
248 =	F8		-8 =	F8

The final result of 0xF8 may be interpreted as 248 for unsigned representation and -8 for a signed representation.

Additionally, $0xF8_{16}$ is the ° (degree symbol) in the ASCII table.

As such, it is very important to have a clear definition of the sizes (byte, halfword, word, etc.) and types (signed, unsigned) of data for the operations being performed.

3.3 Floating-point Representation

The representation issues for floating-point numbers are more complex. There are a series of floating-point representations for various ranges of the value. For simplicity, we will only look primarily at the IEEE 754 32-bit floating-point standard.

3.3.1 IEEE 32-bit Representation

The IEEE 754 32-bit floating-point standard is defined as follows:

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
s	biased exponent								fraction																						

Where *s* is the sign (0 => positive and 1 => negative). When representing floating-point values, the first step is to convert floating-point value into binary.

The following table provides a brief reminder of how binary handles fractional components:

	2^3	2^2	2^1	2^0		2^{-1}	2^{-2}	2^{-3}	
...	8	4	2	1	.	1/2	1/4	1/8	...
	0	**0**	**0**	**0**	**.**	**0**	**0**	**0**	

For example, 100.101_2 would be 4.625_{10}. For repeating decimals, calculating the binary value can be time consuming. However, there is a limit since computers have finite storage.

The next step is to show the value in binary normalized scientific notation. This means that the normalized number should have a single, non-zero leading digit to the left of the decimal point. For example, 8.125_{10} is 1000.001_2 (or 1000.001_2 x 2^0) and in binary normalized scientific notation it would be written as 1.000001 x 2^3 (since the decimal point was moved three places to the left). Of course, if the number was 0.125_{10} the binary would be 0.001_2 (or 0.001_2 x 2^0) and the normalized scientific notation would be 1.0 x 2^{-3} (since the decimal point was moved three places to the right). The numbers after the leading 1, *not* including the leading 1, are stored left-justified in the fraction portion of the word.

The next step is to calculate the *biased exponent,* which is the exponent from the normalized scientific notation plus the defined bias. The bias for the IEEE 754 32-bit floating-point standard is 127_{10}. The result should be converted to a byte (8 bits) and stored in the biased exponent portion of the word.

Note, converting from the IEEE 754 32-bit floating-point representation to the decimal value is done in reverse, however the leading 1 must be added back (as it is not stored in the word). Additionally, the bias is subtracted (instead of added).

3.3.1.1 IEEE 32-bit Representation Examples

This section presents several examples of encoding and decoding floating-point representation for reference.

3.3.1.1.1 Example → -7.75$_{10}$

For example, to find the IEEE 754 32-bit floating-point representation for -7.75$_{10}$:

Example 1: **-7.75**
- determine sign   -7.75 => 1 (since negative)
- convert to binary   -7.75 = -0111.11$_2$
- normalized scientific notation   = 1.1111 x 2^2
- compute biased exponent   2$_{10}$ + 127$_{10}$ = 129$_{10}$
  - and convert to binary   = 10000001$_2$
- write components in binary:
  > sign exponent mantissa
  > 1 10000001 11110000000000000000000
- convert to hex (split into groups of 4)
  > 11000000111110000000000000000000
  > 1100 0000 1111 1000 0000 0000 0000 0000
  >   C   0   F   8   0   0   0   0
- final result: **C0F8 0000$_{16}$**

3.3.1.1.2 Example → -0.125$_{10}$

For example, to find the IEEE 754 32-bit floating-point representation for -0.125$_{10}$:

Example 2: **-0.125**
- determine sign   -0.125 => 1 (since negative)
- convert to binary   -0.125 = -0.001$_2$
- normalized scientific notation   = 1.0 x 2^{-3}
- compute biased exponent   -3$_{10}$ + 127$_{10}$ = 124$_{10}$
  - and convert to binary   = 01111100$_2$
- write components in binary:
  > sign exponent mantissa
  > 1 01111100 00000000000000000000000
- convert to hex (split into groups of 4)
  > 10111110000000000000000000000000
  > 1011 1110 0000 0000 0000 0000 0000 0000
  >   B   E   0   0   0   0   0   0
- final result: **BE00 0000$_{16}$**

3.3.1.1.3 Example → 41440000_{16}

For example, given the IEEE 754 32-bit floating-point representation 41440000_{16} find the decimal value:

> **Example 3:** 41440000_{16}
> - convert to binary
> $$0100\ 0001\ 0100\ 0100\ 0000\ 0000\ 0000\ 0000_2$$
> - split into components
> $$0\ 10000010\ 10001000000000000000000_2$$
> - determine exponent $10000010_2 = 130_{10}$
>   - and remove bias $130_{10} - 127_{10} = 3_{10}$
> - determine sign $0 \Rightarrow$ positive
> - write result $+1.10001 \times 2^3 = +1100.01 = \mathbf{+12.25}$

3.3.2 IEEE 64-bit Representation

The IEEE 754 64-bit floating-point standard is defined as follows:

63	62	52	51	0
s	biased exponent		fraction	

The representation process is the same, however the format allows for an 11-bit biased exponent (which support large and smaller values). The 11-bit biased exponent uses a bias of 1023.

4.0   QtSpim Program Formats

The QtSpim MIPS simulator will be used for programs in this text. The QtSpim simulator has a number of features and requirements for writing MIPS assembly language programs. This includes a properly formatted assembly source file.

A properly formatted assembly source file consists of two main parts; the data section (where data is placed) and the text section (where code is placed). The following sections summarize the formatting requirements and explain each of these sections.

4.1 Assembly Process

The QtSpim effectively assembles the program during the load process. Any major errors in the program format or the instructions will be noted immediately. Assembler errors must be resolved before the program can be successfully executed. Refer to Appendix B regarding the use of QtSpim to load and execute programs.

4.2 Comments

The "#" character represents a comment line. Anything typed after the "#" is considered a comment. Blank lines are accepted.

4.3 Assembler Directives

An assembler directive is a message to the assembler, or the QtSpim simulator, that tells the assembler something it needs to know in order to carry out the assembly process. This includes noting where the data is declared or the code is defined. Assembler directives are *not* executable statements.

Assembler directives start with a ".". Assembler directives are required to define the start and end of data declarations and to define the start and end of procedures/functions. For example, ".data" or ".text".

Additionally, directives are used to declare data. The following sections provide some examples of data declarations using the directives.

4.4 Data Declarations

The data must be declared in the ".data" section. All variables and constants are placed in this section. Variable names must start with a letter followed by letters or numbers (including some special characters such as the "_"), and terminated with a ":" (colon). Variable definitions must include the name, the data type, and the initial value for the variable.   In the definition, the variable name must be terminated with a ":".

The data type must be preceded with a "." (period). The general format is:

```
<variableName>:        .<dataType>      <initialValue>
```

Refer to the following sections for a series of examples using various data types.

The supported data types are as follows:

Declaration	
.byte	8-bit variable(s)
.half	16-bit variable(s)
.word	32-bit variable(s)
.ascii	ASCII string
.asciiz	NULL terminated ASCII string
.float	32 bit IEEE floating-point number
.double	64 bit IEEE floating-point number
.space <n>	<n> bytes of uninitialized memory

These are the primary assembler directives for data declaration. Other directives are referenced in different sections.

4.4.1 Integer Data Declarations

Integer values are defined with the *.word, .half*, or *.byte* directives. Two's complement is used for the representation of negative values. For more information regarding two's complement, refer to the Data Representation section.

The following declarations are used to define the integer variables "wVar1" and "wVar2" as 32-bit word values and initialize them to 500,000 and -100,000.

```
wVar1:      .word       500000
wVar2:      .word       -100000
```

The following declarations are used to define the integer variables "hVar1" and "hVar2" as 16-bit word values and initialize them to 5,000 and -3,000.

```
hVar1:      .half       5000
hVar2:      .half       -3000
```

The following declarations are used to define the integer variables "bVar1" and "bVar2" as 8-bit word values and initialize them to 5 and -3.

```
bVar1:      .byte       5
bVar2:      .byte       -3
```

If a variable is initialized to a value that can not be stored in the allocated space, an assembler error will be generated. For example, attempting to set a byte variable to 500 would be illegal and generate an error.

4.4.2 String Data Declarations

At the assembly level, a string is a series of sequentially defined byte-sized characters, typically terminated with a NULL byte (0x00).

Strings are defined with *.ascii* or *.asciiz* directives. Characters are represented using standard ASCII characters. Refer to Appendix D for a copy of the ASCII table for reference.

The C/C++ style new line, "\n", and tab, "\t" tab are supported within strings.

The following declarations are used to define a string "message" and initialize it to "Hello World".

```
message:    .asciiz     "Hello World\n"
```

In this example, the string is defined as NULL terminated (i.e., after the new line). The NULL is a non-printable ASCII character and is used to mark the end of the string. The NULL termination is standard and is required by the print string system service (to work correctly).

To define a string with multiple lines, the NULL termination would only be required on the final or last line. For example:

```
message:    .ascii      "Line 1: Goodbye World\n"
            .ascii      "Line 2: So, long and thanks "
            .ascii      "for all the fish.\n"
            .asciiz     "Line 3: Game Over.\n"
```

When printed, using the starting address of 'message', everything up-to (but not including) the NULL will be displayed. As such, the declaration using multiple lines is not relevant to the final displayed output.

4.4.3 Floating-Point Data Declarations

Floating-point values are defined with the *.float* (32-bit) or *.double* (64-bit) directives. The IEEE floating-point format is used for the internal representation of floating-point values.

The following declarations are used to define the floating-point variables "pi" a 32-bit floating-point value initialized to 3.14159 and "tao" a 64-bit floating-point values initialized them to 6.28318.

```
pi:         .float      3.14159
tao:        .double     6.28318
```

For more information regarding the IEEE format, refer to the Data Representation section.

4.5 Constants

Constant names must start with a letter, followed by letters or numbers including some special characters such as the "_" (underscore). Constant definitions are created with an "=" sign.

For example, to create some constants named TRUE and FALSE and set them to 1 and 0 respectively:

```
TRUE = 1
FALSE = 0
```

Constants are also defined in the data section. The use of all capitals for a constant is a convention and not required by the QtSpim program. The convention helps programmers more easily distinguish between variables (which can change values) and

constants (which can not change values). Additionally, in assembly language constants are not typed (i.e., not predefined to be a specific size such as 8-bits, 16-bits, 32-bits, or 64-bits).

4.6 Program Code

The code must be preceded by the ".text" directive.

In addition, there are some basic requirements for naming a "main" procedure (i.e., the first procedure to be executed). The ".globl name" and ".ent name" directives are used to define the name of the initial or main procedure. The ".ent" is optional for the QtSpim simulator. *Note,* the **globl** spelled incorrectly is the correct directive. Also, the main procedure must start with a label with the procedure name. The main procedure (as all procedures) should be terminated with the ".end <name>" directive.

In the program template, the <name> would be the name of the main function/procedure, which is "main".

4.7 Labels

Labels are code locations, typically used as a function/procedure name or as the target of a jump. The first use of a label is the main program starting location, which must be named '**main**' which is a specific requirement for the QtSpim simulator.

The rules for a label are as follows:

- Must start with a letter
- May be followed by letters, numbers, or an "_" (underscore).
- Must be terminated with a ":" (colon).
- May only be defined once.

Some examples of a label include:

```
main:
exitProgram:
```

Characters in a label are case-sensitive. As such, `Loop:` and `loop:` are different labels. This can be very confusing initially, so caution is advised.

4.8 Program Template

The following is a very basic template for QtSpim MIPS programs. This general template will be used for all programs.

```
#  Name and general description of program

# ----------------------------------------
#  Data declarations go in this section.

.data
#    program specific data declarations

# ----------------------------------------
#  Program code goes in this section.

.text
.globl     main
.ent       main
main:

# -----
#    your program code goes here.



# -----
#  Done, terminate program.

    li    $v0, 10
    syscall                      # all done!
.end main
```

The initial header (".text", ".globl main", ".ent main", and "main:") will be the same for all QtSpim programs. The final instructions ("li $v0, 10" and "syscall") terminate the program.

A more complete example, with working code, can be found in Appendix A.

5.0   Instruction Set Overview

In assembly-language, instructions are how work is accomplished. In assembly the instructions are simple, single operation commands. In a high-level language, one line might be translated into a series of instructions in assembly-language.

This chapter presents a summary of the basic, most common instructions. The **MIPS Instruction Set** Appendix presents a more comprehensive list of the available instructions.

5.1 Pseudo-Instructions vs Bare-Instructions

As part of the MIPS architecture, the assembly language includes a number of pseudo-instructions. A bare-instruction is an instruction that is directly executed by the CPU. A pseudo-instruction is an instruction that the assembler, or simulator, will recognize but then convert into one or more bare-instructions. This text will focus primarily on the pseudo-instructions.

5.2 Notational Conventions

This section summarizes the notation used within this text which is fairly common in the technical literature. In general, an instruction will consist of the instruction or operation itself (i.e., add, sub, mul, etc.) and the **operands**. The operands refer to where the data (to be operated on) is coming from, or where the result is to be placed.

The following table summarizes the notational conventions used in the remainder of the document.

Operand Notation	Description
Rdest	Destination operand. Must be an integer register. Since it is a destination operand, the contents will be over written with the new result.
Rsrc	Source operand. Must be an integer register. Register value is unchanged after the instruction.

Src	Source operand. Must be an integer register or an integer immediate value. Value is unchanged after the instruction.
FRdest	Destination operand. Must be a floating-point register. Since it is a destination operand, the contents will be overwritten with the new result.
FRsrc	Source operand. Must be a floating-point register. Register value is unchanged after the instruction.
Imm	Immediate value.
Mem	Memory location. May be a variable name or an indirect reference (i.e., a memory address).

By default, the immediate values are decimal or base-10. Hexadecimal or base-16 immediate values may be used but must be preceded with a **0x** to indicate the value is hex. For example, 15_{10} could be entered in hex as **0x0F**.

Refer to the chapter on Addressing Modes for more information regarding memory locations and indirection.

5.3 Data Movement

CPU computations are typically performed using registers. As such, before computations can be performed, data is typically moved into registers from variables (i.e., memory) and when the computations are completed the data would be moved out of registers into other variables.

5.3.1 Load and Store

To support the loading of data from memory (e.g., variables or arrays) into registers and storing of data in register back to memory, there are a series of load and store instructions. The load and store instructions only move data between register and memory. Another instruction is used to move data between registers (as described in the next section).

There are no load or store instructions that will move a value from a memory location directly to another memory location.

The general forms of the load and store instructions are as follows:

Instruction	Description
`l<type>    Rdest, mem`	Load value from memory location into destination register.
`li         Rdest, imm`	Load specified immediate value into destination register.
`la         Rdest, mem`	Load address of memory location into destination register.
`s<type>    Rsrc, mem`	Store contents of source register into memory location.

Assuming the following data declarations:

```
num:        .word       0
wnum:       .word       42
hnum:       .half       73
bnum:       .byte       7
wans:       .word       0
hans:       .half       0
bans:       .byte       0
```

To perform, the basic operations of:

```
num = 27
wans = wnum
hans = hnum
bans = bnum
```

The following instructions could be used:

```
li    $t0, 27
sw    $t0, num              # num = 27
lw    $t0, wnum
sw    $t0, wans             # wans = wnum
lh    $t1, hnum
sh    $t1, hans             # hans = hnum
lb    $t2, bnum
sb    $t2, bans             # bans = bnum
```

For the halfword and byte instructions, only the lower 16-bits or the lower 8-bits are used.

5.3.2 Move

The various forms of the move instructions are used to move data between registers. Both operands must be registers. The most basic move instruction, move, copies the contents of an integer register into another integer register. Another set of move instructions are used to move the contents of registers into or out of the special registers, **$hi** and **$lo**.

In addition, different move instructions are required to move values between integer registers and floating-point registers (as discussed on the floating-point section).

There is no move instruction that will move a value from a memory location directly to another memory location.

The general forms of the move instructions are as follows:

Instruction	Description
`move     Rdest, RSrc`	Copy contents of integer source register into integer destination register.
`mfhi     Rdest`	Copy the contents from the **$hi** register into Rdest register.
`mflo     Rdest`	Copy the contents from the **$lo** register into Rdest register.
`mthi     Rdest`	Copy the contents to the **$hi** register from the Rdest register.
`mtlo     Rdest`	Copy the contents to the **$lo** register from the Rdest register.

For example, the following instructions:

```
li   $t0, 42
move $t1, $t0
```

will move the contents of register **$t0**, 42 in this example, into the **$t1** register.

The *mfhi, mflo, mtho,* and *mtlo* instructions are required only when performing 64-bit integer multiply and divide operations.

The floating-point section will include examples for moving data between integer and floating-point registers.

5.4 Integer Arithmetic Operations

The arithmetic operations include addition, subtraction, multiplication, division, remainder (remainder after division), logical AND, and logical OR. The general format for these basic instructions is as follows:

Instruction	Description
`add     Rdest, Rsrc, Src`	Signed addition Rdest = Rsrc + Src or Imm
`addu    Rdest, Rsrc, Src`	Unsigned addition Rdest = Rsrc + Src or Imm
`sub     Rdest, Rsrc, Src`	Signed subtraction Rdest = Rsrc – Src or Imm
`subu    Rdest, Rsrc, Src`	Unsigned subtraction Rdest = Rsrc – Src or Imm
`mul     Rdest, Rsrc, Src`	Signed multiply with no overflow Rdest = Rsrc * Src or Imm
`mulo    Rdest, Rsrc, Src`	Signed multiply with overflow Rdest = Rsrc * Src or Imm
`mulou   Rdest, Rsrc, Src`	Unsigned multiply with overflow Rdest = Rsrc * Src or Imm
`mult    Rsrc1, Rsrc2`	Signed 64-bit multiply **$hi/$lo** = Rsrc1 * Rsrc2
`multu   Rsrc1, Rsrc2`	Unsigned 64-bit multiply **$hi/$lo** = Rsrc1 * Rsrc2
`div     Rdest, Rsrc, Src`	Signed divide Rdest = Rsrc / Src or Imm

divu	**Rdest, Rsrc, Src**	Unsigned divide Rdest = Rsrc / Src or Imm	
div	**Rsrc1, RSrc2**	Signed divide with remainder **$lo** = Rsrc1 / RSrc2 **$hi** = Rsrc1 % RSrc2	
divu	**Rsrc1, RSrc2**	Unsigned divide with remainder **$lo** = Rsrc1 / RSrc2 **$hi** = Rsrc1 % RSrc2	
rem	**Rdest, Rsrc, Src**	Signed remainder Rdest = Rsrc % Src or Imm	
remu	**Rdest, Rsrc, Src**	Unsigned remainder Rdest = Rsrc % Src or Imm	
abs	**Rdest, Rsrc**	Absolute value Rdest = | Rsrc |	
neg	**Rdest, Rsrc**	Signed negation Rdest = - Rsrc	

These instructions operate on 32-bit registers (even if byte or halfword values are placed in the registers).

Assuming the following data declarations:

```
wnum1:      .word       651
wnum2:      .word       42
wans1:      .word       0
wans2:      .word       0
wans3:      .word       0
hnum1:      .half       73
hnum2:      .half       15
hans:       .half       0
bnum1:      .byte       7

bnum2:      .byte       9
bans:       .byte       0
```

To perform, the basic operations of:

```
wans1 = wnum1 + wnum2
wans2 = wnum1 * wnum2
wans3 = wnum1 % wnum2
hans  = hnum1 * hnum2
bans  = bnum1 / bnum2
```

The following instructions could be used:

```
lw    $t0, wnum1
lw    $t1, wnum2
add   $t2, $t0, $t1
sw    $t2, wans1              # wans1 = wnum1 + wnum2

lw    $t0, wnum1
lw    $t1, wnum2
mul   $t2, $t0, $t1
sw    $t2, wans2              # wans2 = wnum1 * wnum2

lw    $t0, wnum1
lw    $t1, wnum2
rem   $t2, $t0, $t1
sw    $t2, wans3              # wans = wnum1 % wnum2

lh    $t0, hnum1
lh    $t1, hnum2

mul   $t2, $t0, $t1
sh    $t2, hans              # hans = hnum1 * hnum2

lb    $t0, bnum1
lb    $t1, bnum2
div   $t2, $t0, $t1
sb    $t2, bans              # bans = bnum1 / bnum2
```

For the halfword load or store instructions, only the lower 16-bits are used. For the byte instructions, only the lower 8-bits are used.

5.4.1 Example Program, Integer Arithmetic

The following is an example program to compute the volume and surface area of a rectangular parallelepiped.

The formulas for the volume and surface area are as follows:

$$volume \ = \ aSide * bSide * cSide$$
$$surfaceArea \ = \ 2\big(aSide * bSide \ + \ aSide * cSide \ + \ bSide * cSide\big)$$

This example main initializes the *a*, *b*, and *c* sides to arbitrary integer values.

```
#   Example to compute the volume and surface area
#   of a rectangular parallelepiped.

# ---------------------------------------------------------
#   Data Declarations
.data

aSide:              .word       73
bSide:              .word       14
cSide:              .word       16

volume:             .word       0
surfaceArea:        .word       0

# ---------------------------------------------------------
#   Text/code section

.text
.globl      main
.ent        main
main:

# -----
#   Load variables into registers.

        lw      $t0, aSide
        lw      $t1, bSide
        lw      $t2, cSide
```

```
# ----
#  Find volume of a rectangular parallelpiped.
#     volume = aSide * bSide * cSide

    mul   $t3, $t0, $t1
    mul   $t4, $t3, $t2
    sw    $t4, volume

# -----
#  Find surface area of a rectangular parallelepiped.
#     surfaceArea = 2*(aSide*bSide+aSide*cSide+bSide*cSide)

    mul   $t3, $t0, $t1        # aSide * bSide
    mul   $t4, $t0, $t2        # aSide * cSide
    mul   $t5, $t1, $t2        # bSide * cSide
    add   $t6, $t3, $t4
    add   $t7, $t6, $t5
    mul   $t7, $t7, 2
    sw    $t7, surfaceArea

# -----
#  Done, terminate program.

    li    $v0, 10              # call code for terminate
    syscall                    # system call (terminate)
.end main
```

Refer to the system services section for information on displaying the final results to the console.

5.5 Logical Operations

The logical operations include logical AND, logical OR, shift, and rotate instructions. The general format for these instructions is as follows:

Instruction	Description
and Rdest, Rsrc, Src	Logical AND Rdest = Rsrc & Src or Imm
nor Rdest, Rsrc, Src	Logical NOR Rdest = Rsrc ↓ Src or Imm

not	**Rdest, Rsrc, Src**	Logical NOT Rdest = Rsrc ¬ Src or Imm
or	**Rdest, Rsrc, Src**	Logical OR Rdest = Rsrc | Src or Imm
rol	**Rdest, Rsrc, Src**	Rotate left Rdest = Rsrc rotated left Src or Imm places
ror	**Rdest, Rsrc, Src**	Rotate right Rdest = Rsrc rotated right Src or Imm places
sll	**Rdest, Rsrc, Src**	Shift left logical Rdest = Rsrc shift left logical Src or Imm places
sra	**Rdest, Rsrc, Src**	Shift right arithmetic Rdest = Rsrc shift right arithmetic Src or Imm places
srl	**Rdest, Rsrc, Src**	Shift right logical Rdest = Rsrc shift right logical Src or Imm places
xor	**Rdest, Rsrc, Src**	Logical XOR Rdest = Rsrc ^ Src or Imm

The **&** refers to the logical AND operation, the | refers to the logical OR operation, and the ^ refers to the logical XOR operation as per C/C++ conventions. The ↓ refers to the logical NOR operation and the ¬ refers to the logical NOT operation.

These instructions operate on 32-bit registers (even if byte or halfword values are placed in the registers).

Assuming the following data declarations:

```
wnum1:      .word       0x000000ff
wnum2:      .word       0x0000ff00
wans1:      .word       0
wans2:      .word       0
wans3:      .word       0
```

To perform, the basic operations of:

```
wans1 = wnum1 & wnum2
wans2 = wnum1 | wnum2
wans3 = wnum1 ¬ wnum2
```

The following instructions

```
lw    $t0, wnum1
lw    $t1, wnum2
and   $t2, $t0, $t1
sw    $t2, wans1              # wans1 = wnum1 & wnum2

lw    $t0, wnum1
lw    $t1, wnum2
or    $t2, $t0, $t1
sw    $t2, wans2              # wans2 = wnum1 | wnum2

lw    $t0, wnum1
lw    $t1, wnum2
not   $t2, $t0, $t1
sw    $t2, wans3              # wans3 = wnum1 ¬ wnum2
```
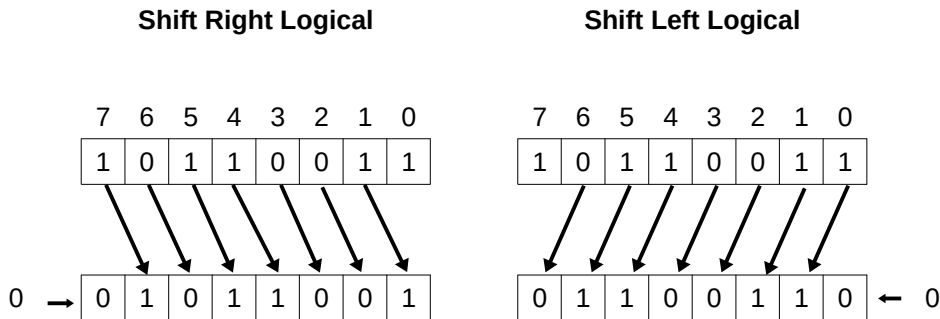
For halfword load or store instructions, only the lower 16-bits are used. For the byte instructions, only the lower 8-bits are used.

5.5.1 Shift Operations

The shift operations shift or move bits within a register. Two typical reasons for shifting bits include isolating a subset of the bits within an operand for some specific purpose or possibly for performing multiplication or division by powers of two. The two shift operations are a logical shift and an arithmetic shift.

5.5.1.1 Logical Shift

The logical shift is a bitwise operation that shifts all the bits of its source register by the specified number of bits and places the result into the destination register. The bits can be shifted left or right as needed. Every bit in the source operand is moved the specified number of bit-positions, and the newly vacant bit-positions are filled in with zeros. The following diagram shows how the right and left shift operations work for byte sized operands.

Shift Right Logical   **Shift Left Logical**

7	6	5	4	3	2	1	0
1	0	1	1	0	0	1	1

0 →

0	1	0	1	1	0	0	1

7	6	5	4	3	2	1	0
1	0	1	1	0	0	1	1

0	1	1	0	0	1	1	0

← 0

The logical shift treats the operand as a sequence of bits rather than as a number.

The shift instructions may be used to perform unsigned integer multiplication and division operations for powers of 2. Powers of two would be 2, 4, 8, etc. up to the limit of the operand size (32-bits for register operands).

In the examples below, 23 is divided by 2 by performing a shift right logical one bit. The resulting 11 is shown in binary. Next, 13 is multiplied by 4 by performing a shift left logical two bits. The resulting 52 is shown in binary.

Shift Right Logical
Unsigned Division

Shift Left Logical
Unsigned Multiplication

0	0	0	1	0	1	1	1	= 23

0	0	0	0	1	0	1	1	= 11

0	0	0	0	1	1	0	1	= 13

0	0	1	1	0	1	0	0	= 52

As can be seen in the examples, a 0 was entered in the newly vacated bit locations on either the right or left (depending on the operation).

5.5.1.2 Arithmetic Shift

The arithmetic shift right is also a bitwise operation. This instruction shifts all bits of the source register by the specified number of bit-positions and places the result into the destination register. Every bit in the source operand is moved the specified number of bit-positions, and the newly vacant bit-positions on the left are filled in. The original leftmost bit (the sign bit) is replicated to fill in all the vacant positions. This is referred to as sign extension. The following diagram shows how the shift right arithmetic operations work for a byte sized operand.

Shift Right Arithmetic

7	6	5	4	3	2	1	0
1	0	1	1	0	0	1	1

1	1	0	1	1	0	0	1

The arithmetic shift treats the operand as a signed number and extends the sign which would be negative in this example.

However, the arithmetic shift rounds up and the standard divide instruction truncates. As such, the arithmetic shift is not typically used to replace the signed divide instruction.

5.5.1.3 Shift Operations, Examples

This section provides a series of examples using the logical shift operations.

Assuming the following data declarations:

```
data1:     .word     0x000000ff
result1:   .word     0
result2:   .word     0
```

To perform, the basic operations of:

```
result1 = wnum1, rotate left 1 bit
result2 = wnum1, rotate right 1 bit
```

The following instructions

```
lw    $t0, wnum1
lw    $t1, wnum2
rol   $t2, $t0, $t1
sw    $t2, wans3        # wans3 = wnum1, rotate left 1 bit

lw    $t0, wnum1
lw    $t1, wnum2
ror   $t2, $t0, $t1
sw    $t2, wans4        # wans3 = wnum1, rotate right 1 bit
```

For halfword instructions, only the lower 16-bits are used. For the byte instructions, only the lower 8-bits are used.

To perform the operation, **value * 8**, it would be possible to shift the number in the variable one bit for each power of two, which would be three bits in this example.

Assuming the following data declarations:

```
value:    .word     17
answer:   .word     0
```

The following instructions could be used to multiply a value by 8.

```
lw    $t0, value
sll   $t1, $t0, 3
sw    $t1, answer           # answer = value * 8
```

The final value in answer would be 17 * 8 or 136.

In the context of an encoded MIPS instruction, the upper 6-bits of a 32-bit word represent the OP or operation field. If a program was analyzing code, it might be desirable to isolate these bits for comparison. One way this can be performed is to use a logical right shift to move the upper six bits into the position of the lower 6-bits.

The instruction:

```
add   $t1, $t1, 1
```

will be translated by the assembler into the hex value of **0x2129001**.

Assuming the following data declarations:

```
inst1:    .word     0x2129001
inst1Op1: .word     0
```

To mask out the OP field (upper 6-bits) for *inst1* and place it in the variable *instOp1* (lower 6-bits), the following instructions could be used:

```
lw     $t0, inst1
srl    $t1, $t0, 26
sw     $t1, instOp1
```

This can be done in one step since the logical shift will insert all 0's into the newly vacated bit locations.

5.6 Control Instructions

Program control refers to basic programming structures for iteration and comparisons such as IF statements and looping. All of the high-level language control structures must be performed with the limited assembly-language control structures. For example, an IF-THEN-ELSE statement does not exist at the assembly-language level. Assembly-language provides an unconditional branch (or jump), and a conditional branch or an IF statement that will jump to a target label or not jump (as per the conditional expression).

The control instructions refer to unconditional and conditional branching. Branching is required for basic conditional statements (i.e., IF statements) and looping.

5.6.1 Unconditional Control Instructions

The unconditional instruction provides an unconditional jump to a specific location.

Instruction	Description
j <label>	Unconditionally branch to the specified label.

The "**b**" (branch) may be used instead of the "**j**" (jump). Both are encoded as the same instruction (an unconditional jump). An error is generated by QtSpim if the label is not defined.

5.6.2 Conditional Control Instructions

The conditional instruction provides a conditional jump based on a comparison. In high-level language terms, this is a basic IF statement.

The conditional control instructions include the standard set; branch equal, branch not equal, branch less than, branch less than or equal, branch greater than, and branch greater than or equal.

The general format for these basic instructions is as follows:

Instruction	Description
`beq <Rsrc>, <Src>, <label>`	Branch to label if <Rsrc> and <Src> are equal
`bne <Rsrc>, <Src>, <label>`	Branch to label if <Rsrc> and <Src> are not equal
`blt <Rsrc>, <Src>, <label>`	Signed branch to label if <Rsrc> is less than <Src>
`ble <Rsrc>, <Src>, <label>`	Signed branch to label if <Rsrc> is less than or equal to <Src>
`bgt <Rsrc>, <Src>, <label>`	Signed branch to label if <Rsrc> is greater than <Src>
`bge <Rsrc>, <Src>, <label>`	Signed branch to label if <Rsrc> is greater than or equal to <Src>
`bltu <Rsrc>, <Src>, <label>`	Unsigned branch to label if <Rsrc> is less than <Src>
`bleu <Rsrc>, <Src>, <label>`	Unsigned branch to label if <Rsrc> is less than or equal to <Src>
`bgtu <Rsrc>, <Src>, <label>`	Unsigned branch to label if <Rsrc> is greater than <Src>
`bgeu <Rsrc>, <Src>, <label>`	Unsigned branch to label if <Rsrc> is greater than or equal to <Src>

These instructions operate on 32-bit registers (even if byte or halfword values are placed in the registers).

In addition, these conditional control instructions can be modified by adding or appending a 'z' to the end which means a comparison to zero (0) without typing the immediate 0 in the instruction.

For example, the following instruction,

```
bne    $t0, 0, loop1
```

could be written as,

```
bnez   $t0, loop1
```

which does exactly the same thing. This short-handed method is used in some of the text examples. A more complete list is included in Appendix C.

5.6.3 Example Program, Sum of Squares

The following is an example program to find the sum of squares from 1 to *n*. For example, the sum of squares from 1 to 10 is as follows:

$$1^2 + 2^2 + \cdots + 10^2 \ = \ 385$$

This example program initializes the *n* to 10 to match the example above example. Other limits can be specified as desired.

```
#  Example program to compute the sum of squares.
# --------------------------------------------------------
#  Data Declarations

.data
n:                  .word       10
sumOfSquares:   .word       0

# --------------------------------------------------------
#  text/code section

.text
.globl      main
.ent        main
main:

# -----
#  Compute sum of squares from 1 to n.

        lw    $t0, n                 #
        li    $t1, 1                 # loop index (1 to n)
        li    $t2, 0                 # sum
```

```
    sumLoop:
         mul   $t3, $t1, $t1          # index^2
         add   $t2, $t2, $t3

         add   $t1, $t1, 1
         ble   $t1, $t0, sumLoop

         sw    $t2, sumOfSquares

    # -----
    #  Done, terminate program.

         li    $v0, 10                 # call code for terminate
         syscall                       # system call
    .end main
```

Refer to the system services section for information on displaying the final results to the console.

5.7 Floating-Point Instructions

This section presents a summary of the basic, most common floating-point arithmetic instructions. The *MIPS Instruction Set* Appendix presents a more comprehensive list of the available instructions.

5.7.1 Floating-Point Register Usage

The floating-point instructions are similar to the integer instructions, however, the floating-point register must be used with the floating-point instructions. Specifically, this means the architecture does not support the use of integer registers for any floating-point arithmetic operations.

When single-precision (32-bit) floating-point operation is performed, the specified 32-bit floating-point register is used. When a double-precision (64-bit) floating-point operation is performed, two 32-bit floating-point registers are used; the specified 32-bit floating-point register and the next numerically sequential register is used by the instruction. For example, a double-precision operation using **$f12** will use automatically **$f12** and **$f13**.

5.7.2 Floating-Point Data Movement

Floating-point CPU computations are typically performed using floating-point registers. As such, before computations can be performed, data is typically moved into the floating-point registers from other floating-point registers or variables (i.e., memory). When a computation is completed the data might be moved out of the floating-point register into a variable or another floating-point register.

To support the loading of data from memory into floating-point registers and storing of data in floating-point registers to memory, there are a series of specialized load and store instructions. The basic format is the same as the integer operations, however the type is either ".s" for single-precision 32-bit IEEE floating-point representation or ".d" for double-precision 64-bit IEEE floating-point representation. More information regarding the representations can be found in Chapter 2, *Data Representation*.

The general forms of the floating-point load and store instructions are as follows:

Instruction	Description
`l.<type>     FRdest, mem`	Load value from memory location memory into destination register.
`s.<type>     FRsrc, mem`	Store contents of source register into memory location.
`mov.<type>  Frdest, FRsrc`	Copy the contents of source register into the destination register.

In this case, the floating-point types are ".s" for single-precision and ".d" for double-precision.

Assuming the following data declarations:

```
fnum1:     .float          3.14
fnum2:     .float          0.0
dnum1:     .double         6.28
dnum2:     .double         0.0
```

The ".float" directive declares a variable as a 32-bit floating-point value and the ".double" declares a variable as a 64-bit floating-point variable.

To perform, the basic operations of:

```
fnum2 = fnum1
dnum2 = dnum1
```

The following instructions :

```
l.s     $f6, fnum1
s.s     $f6, fnum2              # fnum2 = fnum1

l.d     $f6, dnum1
mov.d   $f8, $f6               # unnecessary use of mov
                               # just as an example
s.d     $f8, dnum2            # dnum2 = dnum1
```

The two double-precision operations (l.d and mov.d) reference registers **$f6** and **$f8** but use registers **$f6**/**$f7** and **$f8**/**$f9** to hold each of the two 64-bit values.

5.7.3 Integer / Floating-Point Register Data Movement

The arithmetic instructions require either floating-point registers or integer registers and will not allow a combination. In order to move data between integer and floating-point registers, special instructions are required. As noted in Chapter 2, *MIPS Architecture Overview*, the floating-point operations are performed in a floating-point co-processor.

The general form of the integer and floating-point data movement instructions are as follows:

Instruction	Description
`mfc1       Rdest, FRsrc`	Copy the contents from co-processor 1 (FPU) float register FRsrc into Rdest integer register.
`mfc1.d     Rdest, FRsrc`	Copy the contents from co-processor 1 (FPU) float registers FRsrc and FRsrc+1 into integer registers Rdest and Rdest+1.
`mtc1       Rsrc, FRdest`	Copy the contents from integer Rsrc register to co-processor 1 (FPU) float register FRdest.

`mtc1.d     Rsrc, FRdest`	Copy the contents from integer registers Rdest and Rdest+1 to co-processor 1 (FPU) float registers FRsrc and FRsrc+1.

Note, the above instructions use a 1 (number one) and not a lower-case letter L.

For example, assuming an integer value is in integer register **$s0**, to copy the value into floating-point register **$f12**, the following instruction could be used.

```
mtc1       $s0, $f12
```

To copy the contents of **$f12**, into an integer register **$t1**, the following instruction could be used.

```
mfc1       $t1, $f12
```

The value copied has not been converted into a different representation.

In this example, the integer value in **$s0** that was copied into **$f12** is still represented as an integer in two's complement. As such, the value in **$f12** is not ready for any floating-point arithmetic operations. The representation of the value must be converted (see next section).

5.7.4 Integer / Floating-Point Conversion Instructions

When data is moved between integer and floating-point registers, the data representation must be addressed. For example, when moving an integer value from an integer register into a floating-point register, the data is still represented as an integer value in two's complement. Floating-point operations require an appropriate floating-point representation (32-bit or 64-bit). When data is moved between integer and floating-point registers, a data conversion would typically be required.

The general format for the conversion instructions is as follows:

Instruction	Description
`cvt.d.s    FRdest, FRsrc`	Convert the 32-bit floating-point value in register FRsrc into a double precision value and put it in register FRdest.

Instruction	Description
`cvt.d.w    FRdest, FRsrc`	Convert the 32-bit integer in register FRsrc into a double precision value and put it in register FRdest.
`cvt.s.d    FRdest, FRsrc`	Convert the 64-bit floating-point value in register FRsrc into a 32-bit floating-point value and put it in register FRdest.
`cvt.s.w    FRdest, FRsrc`	Convert the 32-bit integer in register FRsrc into a 32-bit floating-point value and put it in register FRdest.
`cvt.w.d    FRdest, FRsrc`	Convert the 64-bit floating-point value in register FRsrc into a 32-bit integer value and put it in register FRdest.
`cvt.w.s    FRdest, FRsrc`	Convert the 32-bit floating-point value in register FRsrc into a 32-bit integer value and put it in register FRdest.

Assuming the following data declarations:

```
iNum:       .word       42
fNum:       .float      0.0
```

To convert the integer value in variable *iNum* and place it as a 32-bit floating-point value in variable *fNum*, the following instructions could be used:

```
lw          $t0, iNum
mtc1        $t0, $f6
cvt.s.w     $f8, $f6
s.s         $f8, fNum
```

This code fragment loads the integer value in variable *iNum* into **$t0**, and then copies the value into **$f6**. The integer value in **$f6** is converted into a 32-bit floating-point value and placed in **$f8**. The 32-bit floating-point value is then copied into the *fNum* variable. The conversion instruction could have over-written the **$f6** register.

Assuming the following data declarations:

```
pi:          .double    3.14
intPi:       .word      0
```

To convert the 64-bit floating-point value in variable *pi* and place it as a 32-bit integer value in variable *intPi*, the following instructions could be used:

```
l.d          $f10, pi
cvt.w.d      $f12, $f10
mfc1         $t1, $f12
sw           $t1, intPi
```

This code fragment initially loads the 64-bit floating-point value into **$f10**. The 64-bit floating-point value in **$f10** is converted into a 32-bit integer value and placed in **$f12**. The integer value in **$f12** is copied into **$t1** and then copied into the variable *intPi*. Since conversion from floating-point truncates, the final value in *intPi* is 3.

5.7.5 Floating-Point Arithmetic Operations

The arithmetic operations include addition, subtraction, multiplication, division, remainder (remainder after division), logical AND, and logical OR.

The general format for these basic instructions is as follows:

Instruction	Description
`add.<type> FRdest, FRsrc, FRsrc`	FRdest = FRsrc + FRsrc
`sub.<type> FRdest, FRsrc, FRsrc`	FRdest = FRsrc - FRsrc
`mul.<type> FRdest, FRsrc, FRsrc`	FRdest = FRsrc * FRsrc
`div.<type> FRdest, FRsrc, FRsrc`	FRdest = FRsrc / FRsrc
`rem.<type> FRdest, FRsrc, FRsrc`	FRdest = FRsrc % FRsrc

Assuming the following data declarations:

```
fnum1:       .float     6.28318
fnum2:       .float     3.14159
fans1:       .float     0.0
fans2:       .float     0.0
```

```
dnum1:      .double     42.3
dnum2:      .double     73.6
dans1:      .double     0.0
dans2:      .double     0.0
```

To perform, the basic operations of:

```
fans1 = fnum1 + fnum2
fans2 = fnum1 * fnum2
dans1 = dnum1 - dnum2
dans2 = dnum1 / dnum2
```

The following instructions:

```
l.s         $f4, fnum1
l.s         $f6, fnum2
add.s       $f8, $f4, $f6
s.s         $f8, fans1              # fans1 = fnum1 + fnum2

mul.s       $f10, $f4, $f6
s.s         $f10, fans2            # fans2 = fnum1 * fnum2

l.d         $f4, dnum1
l.d         $f6, dnum2
sub.d       $f8, $f4, $f6
s.d         $f8, dans1             # dans1 = dnum1 - dnum2

div.d       $f10, $f4, $f6
s.d         $f10, dans2            # dans2 = dnum1 / dnum2
```

For the double-precision instructions, the specified register and the next numerically sequential register is used. For example, the **l.d** instruction sets the **$f4** and **$f5** 32-bit registers with the 64-bit value.

5.7.6 Example Programs

This section provides some example using the floating-point instructions to perform some basic calculations.

5.7.6.1 Example Program, Floating-Point Arithmetic

The following is an example program to compute the surface area and volume of a sphere.

The formulas for the surface area and volume of a sphere are as follows:

$$surfaceArea \; = \; 4.0 * pi * radius^2$$

$$volume \; = \; \frac{4.0 * pi}{3.0} * radius^3$$

This example main initializes the **radius** to an arbitrary floating-point value.

```
#   Example program to calculate the surface area
#   and volume of a sphere given the radius.
# --------------------------------------------------------
#   Data Declarations

.data
pi:             .float      3.14159
fourPtZero:     .float      4.0
threePtZero:    .float      3.0
radius:         .float      17.25
surfaceArea:    .float      0.0
volume:         .float      0.0


# --------------------------------------------------------
#   text/code section

.text
.globl      main
.ent        main
main:

#   Compute: (4.0*pi) which is used for both equations.

        l.s         $f2, fourPtZero
        l.s         $f4, pi
        mul.s       $f4, $f2, $f4       # 4.0 * pi
```

```
        l.s         $f6, radius             # radius

    # -----
    #  Calculate surface area of a sphere.
    #     surfaceArea = 4.0 * pi * radius^2

        mul.s       $f8, $f6, $f6           # radius^2
        mul.s       $f8, $f4, $f8           # 4.0*pi * radius^2
        s.s         $f8, surfaceArea        # store final answer

    # -----
    #  Calculate volume of a sphere.
    #     volume = (4.0 * pi / 3.0) * radius^3

        l.s         $f8, threePtZero

        div.s       $f2, $f4, $f8           # (4.0 * pi / 3.0)

        mul.s       $f10, $f2, $f2
        mul.s       $f10, $f10, $f6         # radius^3

        mul.s       $f12, $f6, $f10         # * 4.0*pi/3.0

        s.s         $f12, volume            # store final answer

    # -----
    #  Done, terminate program.

        li          $v0, 10                 # terminate
        syscall                             # system call
    .end main
```

Refer to the system services section for information on displaying the final results to the console.

5.7.6.2 Example Program, Integer / Floating-Point Conversion

The following is an example program to sum an array of integer values and compute the average as a floating-point value. This requires conversion of 32-bit integer values into 32-bit floating-point values.

```
#   Example program to sum an array of integers
#   and compute the float average.

# ---------------------------------------------------------
#   Data Declarations

.data
iArray:      .word       1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12
length:      .word       12

iSum:        .word       0
fAve:        .float      0.0

# ---------------------------------------------------------
#   Text/code section

.text
.globl      main
.ent        main
main:

# -----
#   Find the sum of the integer numbers.

      la    $t0, iArray             # array starting addr
      lw    $t1, length            # array length
      li    $t2, 0                 # set sum=0

sumLoop:
      lw          $t3, ($t0)            # get iArray(n)
      add         $t2, $t2, $t3         # sum=sum+iArray(n)
      addu        $t0, $t0, 4           # update iArray addr
      sub         $t1, $t1, 1
      bnez        $t1, sumLoop

      sw          $t2, iSum             # save integer sum

      mtc1        $t2, $f6              # move to flt reg
      cvt.s.w     $f6, $f6              # cvt to flt format

      lw          $t1, length
      mtc1        $t1, $f8              # move to float reg
      cvt.s.w     $f8, $f8              # cvt to float format
```

```
        div.s       $f10, $f6, $f8          # sum / length
        s.s         $f10, fAve

    # -----
    #   Done, terminate program.

        li    $v0, 10                         # terminate
        syscall                               # system call
    .end main
```

6.0   Addressing Modes

This chapter provides basic information regarding addressing modes and the associated address manipulations on the MIPS architecture. The addressing modes are the supported methods for specifying the value or address of a data item being accessed (read or written). This might include an actual value, the name of a variable, or the location in an array.

Since the MIPS architecture, as simulated in the QtSpim simulator, is a 32-bit architecture, all addresses are words (32-bits).

6.1 Direct Mode

Direct addressing mode is when the register or memory location contains the actual values.

For example:

```
lw    $t0, var1
lw    $t1, var2
```

Registers and variables **$t0**, **$t1**, *var1*, and *var2* are all accessed in direct mode addressing.

6.2 Immediate Mode

Immediate addressing mode is when the actual value is one of the operands.

For example:

```
li    $t0, 57
add   $t0, $t0, 57
```

The value 57 is immediate mode addressing. The register **$t0** is direct mode addressing.

6.3 Indirection

The pair of parenthesis, **( )**'s, are used to denote an indirect memory access. An indirect memory access means the CPU will read the provided address and then go to that address to access the value located there. This involves more work for the CPU than the previously presented addressing modes (direct and immediate). This is typically how elements are accessed in a list or array. For example, to get a value from a list of longs:

```
la    $t0, lst
lw    $s1, ($t0)
```

The address, in **$t0**, is a word size (32-bits). Memory is byte addressable. As such, if the data items in "lst" (from above) are words, then four must be added to get the next element.

For example, the instructions:

```
add   $t0, $t0, 4
lw    $s2, ($t0)
```

will get the next word value in array (named *lst* in this example).

A form of displacement addressing is allowed. For example, to get the second item from a list of word sized values:

```
la    $t0, lst
lw    $s1, 4($t0)
```

The "4" is added to the address before the memory access. However, the register is not changed. Thus, the location or address being accessed is displaced or temporarily changed as needed.

6.3.1 Bounds Checking

In a high-level language, the compiler is capable of ensuring that the index for an element in an array is legal and within the boundary of the array being accessed. Thus, the compiler can issue an error message and help identify when and where a program is trying to access beyond the end of an array (e.g., accessing the 110th element of a 100 element array).

This type of bounds checking is not available at the assembly-language level.

If the assembly-language program attempts to access the 110[th] element of an array, the value at that memory location will be returned with no error. Of course, the value returned is not likely to be useful.

If the memory access attempting to be accessed is outside the general scope of the program, an exception will be generated. An exception is a run-time error, and the QtSpim simulator will provide the line where the exception occurred. For example, attempting to access a memory location in the reserved section would not be allowed and thus generate an exception. This could easily occur if the programmer uses a register with a data item instead of a correct address.

Additionally, no error is generated when a program attempts to access a word (32-bits) in an array of halfwords (16-bits). In this case two halfwords will be read into the registers and treated as a single value. Of course, the value will not be correct or useful.

6.4 Examples

This section provides some example using the addressing modes to access arrays and perform basic calculations.

6.4.1 Example Program, Sum and Average

The following example computes the sum and average for an array integer values. The values are calculated and saved into memory variables.

```
#  Example to compute the sum and integer average
#  for an array of integer values.
# ----------------------------------------------------
#  Data Declarations

.data

array:      .word   1,  3,  5,  7,  9, 11, 13, 15, 17, 19
            .word  21, 23, 25, 27, 29, 31, 33, 35, 37, 39
            .word  41, 43, 45, 47, 49, 51, 53, 55, 57, 59
length:     .word  30

sum:        .word   0
average:    .word   0
```

```
# --------------------------------------------------------
#  Basic approach:
#     - loop through the array
#           access each value, update sum
#     - calculate the average

.text
.globl      main
.ent        main
main:

# -----
#  Loop through the array to calculate sum

      la    $t0, array              # array starting address
      li    $t1, 0                  # loop index, i=0
      lw    $t2, length             # length
      li    $t3, 0                  # initialize sum=0

sumLoop:
      lw    $t4, ($t0)              # get array[i]
      add   $t3, $t3, $t4           # sum = sum + array[i]

      add   $t1, $t1, 1             # i = i+1
      add   $t0, $t0, 4             # update array address

      blt   $t1, $t2, sumLoop       # if i<length, continue

      sw    $t3, sum                # save sum

# -----
#  Calculate average
#     note, sum and length set in section above.

      div  $t5, $t3, $t2            # ave = sum / length
      sw    $t5, average

# -----
#  Done, terminate program.

      li    $v0, 10                 # terminate
      syscall                       # system call
.end main
```

This example program does not display the results to the screen. For information regarding displaying values and strings to output (console), refer to the QtSpim System Services section.

6.4.2 Example Program, Median

The following example finds the median for a sorted array of values. In this example, the length is given as always even. As such, the integer median is the integer average for the two middle values. Specifically, the formula for median is:

$$medianEvenOnly \;=\; \frac{\big(array[length/2] + array[length/2-1]\big)}{2}$$

The 'length/2' notation refers to using division by two to generate the correct index of the appropriate value from the array. In assembly, we must convert the index into the offset from the base address (i.e., starting address) of the array. Since the array elements in this example are words (i.e., 4 bytes), it will be necessary to multiply by four to convert the index into an offset. That offset is from the start of the array, so the final address is the array base address plus the offset.

This requires a series of calculations as demonstrated in the following example.

```
#  Example to find the median of a sorted
#  array of integer values of even length.
# -------------------------------------------------------
#  Data Declarations

.data
array:      .word    1,  3,  5,  7,  9, 11, 13, 15, 17, 19
            .word   21, 23, 25, 27, 29, 31, 33, 35, 37, 39
            .word   41, 43, 45, 47, 49, 51, 53, 55, 57, 59

length:     .word   30
median:     .word    0

# -------------------------------------------------------
#  text/code section
#  The median for an even length array is defined as:
#     median = ( array[len/2] + array[len/2-1] ) / 2
#  Note, the len/2 is the index. Must convert the index
#  into the an offset from the base address (of the
#  array.  Since the data is words (4 bytes), multiply
```

```
    #  the index by four to convert to the offset.

    .text
    .globl    main
    .ent      main
    main:
        la    $t0, array          # starting addr of array
        lw    $t1, length         # value of length

        div   $t2, $t1, 2         # length / 2
        mul   $t3, $t2, 4         # cvt index into offset
        add   $t4, $t0, $t3       # add base addr of array

        lw    $t5, ($t4)          # get array[len/2]
        sub   $t4, $t4, 4         # addr of prev value

        lw    $t6, ($t4)          # get array[len/2-1]

        add   $t7, $t6, $t5       # a[len/2] + a[len/2-1]
        div   $t8, $t7, 2         #  / 2

        sw    $t8, median         # save median

    # -----
    #  Done, terminate program.

        li    $v0, 10             # terminate
        syscall                   # system call
    .end main
```

This example program does not display the results to the screen. For information regarding displaying values and strings to output (console), refer to the QtSpim System Services section.

Finding the median for an odd length list is left to the reader as an exercise.

7.0   Stack

In a computer, a stack is a type of data structure where items are added and then removed from the stack in reverse order. That is, the most recently added item is the very first one that is removed. This is often referred to as Last-In, First-Out (LIFO).

A stack is heavily used in programming for the storage of information during procedure or function calls. The following chapter provides information and examples regarding procedure and function calls.

Adding an item to a stack is referred to as a *push* or push operation. Removing an item from a stack is referred to as a *pop* or pop operation.

It is generally expected that the reader will be familiar with the general concept of a stack.

7.1 Stack Example

To demonstrate the usage of the stack, given an array, `a = {7, 19, 37}`, consider the operations:

```
push  a[0]
push  a[1]
push  a[2]
```

Followed by the operations:

```
pop   a[0]
pop   a[1]
pop   a[2]
```

The initial push will push the 7, followed by the 19, and finally the 37. Since the stack is last-in, first-out, the first item popped off the stack will be the last item pushed, or 37 in this example. The 37 is placed in the first element of the array (over-writing the 7). As this continues, the order of the array elements is reversed.

The following diagram shows the progress and the results.

stack	stack	stack	stack	stack	stack
		37			
	19	19	19		
7	7	7	7	7	empty

push a[0]	push a[1]	push a[2]	pop a[0]	pop a[1]	pop a[2]
a = {7, 19, 37}	a = {7, 19, 37}	a = {7, 19, 37}	a = {37, 19, 37}	a = {37, 19, 37}	a = {37, 19, 7}
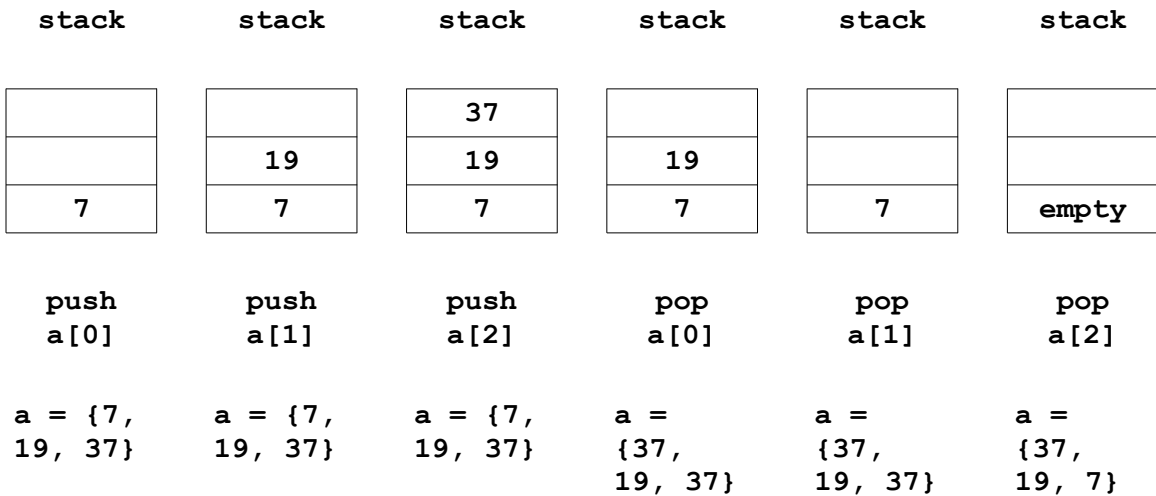
The following sections provide more detail regarding the implementation and applicable instructions.

7.2 Stack Implementation

The current top of the stack is pointed to by the **$sp** register. The stack grows downward in memory and it is generally expected that all items pushed and/or popped should be of word size (32-bit).

There is no push or pop instruction. Instead, you must perform the push and pop operations manually.

While it is possible to push/pop items of various sizes (byte, halfword, etc.) it is not recommended. For such operations, it is recommended to use the entire word (4-bytes).

7.3 Push

For example, a push would subtract the **$sp** by 4 bytes and then copy the operand to that location (in that order). The instructions to push **$t9** would be implemented as follows:

```
    subu    $sp, $sp, 4
    sw      $t9, ($sp)
```

Which will place the contents of the **$t9** register at the top of the stack.

7.4 Pop

A pop would copy the top of the stack to the operand and then add 4 bytes (in that order). To pop the stack into **$t2**, the instructions would be as follows:

```
lw     $t2, ($sp)
addu   $sp, $sp, 4
```

Which will copy the contents of the top of the stack into the **$t2** register.

7.5 Multiple push's/pop's

The preferred method of performing multiple pushes or pops is to perform the **$sp** adjustment only once. For example, to push registers, **$s0**, **$s1**, and **$s2**:

```
subu   $sp, $sp, 12
sw     $s0, ($sp)
sw     $s1, 4($sp)
sw     $s2, 8($sp)
```

And, the commands to pop registers, **$s0**, **$s1**, and **$s2** as follows:

```
lw     $s0, ($sp)
lw     $s1, 4($sp)
lw     $s2, 8($sp)
addu   $sp, $sp, 12
```

By performing the stack adjustment only once, it is more efficient for the architecture to execute.

7.6 Example Program, Stack Usage

The following example uses a stack to reverse the elements in an array. The program will push all elements of the array to the stack and then pop all elements back into the array. This will place the elements back into the array in reverse order based on the basic functionality of the stack.

```
#  Example to reverse values in an array
#  by using the stack.
# -------------------------------------------------------
#  Data Declarations
```

```
    .data

array:      .word      1,  3,  5,  7,  9, 11, 13, 15, 17, 19
            .word     21, 23, 25, 27, 29, 31, 33, 35, 37, 39
            .word     41, 43, 45, 47, 49, 51, 53, 55, 57, 59
length:     .word     30

# --------------------------------------------------------
#  Text/code section

#  Basic approach:
#     - loop to push each element onto the stack
#     - loop to pop each element off the stack
#  Final result is all elements reversed.

.text
.globl      main
.ent        main
main:

# -----
#  Loop to read items from array and push to stack.

        la   $t0, array           # array starting address
        li   $t1, 0               # loop index, i=0
        lw   $t2, length          # length

pushLoop:
        lw   $t4, ($t0)           # get array[i]

        subu $sp, $sp, 4          # push array[i]
        sw   $t4, ($sp)

        add  $t1, $t1, 1          # i = i+1
        add  $t0, $t0, 4          # update array address

        blt  $t1, $t2, pushLoop   # if i<length, continue

# -----
#  Loop to pop items from stack and write into array.

        la   $t0, array           # array starting address
        li   $t1, 0               # loop index, i=0
```

```
        lw    $t2, length              # length (redundant line)

popLoop:
        lw    $t4, ($sp)
        addu  $sp, $sp, 4              # pop array[i]

        sw    $t4, ($t0)              # set array[i]

        add   $t1, $t1, 1             # i = i+1
        add   $t0, $t0, 4             # update array address

        blt   $t1, $t2, popLoop       # if i<length, continue

# -----
#  Done, terminate program.

        li    $v0, 10                 # terminate
        syscall                       # system call
.end main
```

It must be noted that there are easier ways to reverse a set of numbers, but they would not help demonstrate stack operations.

8.0   Procedures/Functions

This chapter provides an overview of using assembly language procedures/functions. In C/C++ a procedure is referred to as a void function. Other languages refer to such functions as procedures. A function returns a single value in a more mathematical sense. C/C++ refers to functions as value returning functions.

With regard to calling a procedure/function, there are two primary activities; linkage and argument transmission. Each is explained in the following sections. Additionally, using procedures/functions in MIPS assembly language requires the use of a series of special purpose registers. These special purpose registers are part of the basic integer register set but have a dedicated purpose based upon standardized and conventional usage.

8.1 MIPS Calling Conventions

When writing MIPS assembly-language procedures, the MIPS standard calling conventions should be utilized. This ensures that the code can be more effectively re-used, can interact with other compiler-generated code or mixed-language programs, and utilize high-level language libraries.

The calling conventions address register usage, argument passing and register preservation.

There are two categories of procedures as follows:

- Non-leaf procedures
  ○ These procedures call other procedures.
- Leaf procedures
  ○ These procedures do not call other procedures (or themselves).

The standard calling convention specifies actions for the *caller* (routine that is calling) and the *callee* (routine that is being called). The specific requirements for each are detailed in the following sections.

8.2 Procedure/Function Format

The basic format for a procedure/function declaration uses a global declaration directive
(".globl <procName>"), an entry point directive (".ent <procName>"), and an entry label
for the procedure. Generally, a procedure declaration is terminated with an end
directive (".end <procName>"). The general syntax is as follows:

```
.globl  functionName
.ent    functionName
functionName:

# code goes here

.end  functionName
```

The use of the ".end <functionName>" directive is optional in the QtSpim simulator.

8.3 Caller Conventions

The calling convention addresses specific requirements for the *caller* or routine that is
calling a procedure.

- The calling procedures are expected to save any non-preserved registers (**$a0** -
  $a3, **$t0** - **$t9**, **$v0**, **$v1**, **$f0** - **$f10** and **$f16** - **$f18**) that are required after the
  call is completed.
- The calling procedure should pass all arguments.
  - The first argument is passed in either **$a0** or **$f12** (**$a0** if integer or **$f12** if
    float single or double precision).
  - The second argument is passed in either **$a1** or **$f14** (**$a1** if integer or **$f14** if
    float single or double precision).
  - The third argument is passed in **$a2** (integer only).
  - If the third argument is float, it must be passed on the stack.
  - The fourth argument is passed in **$a3** (integer only).
  - If the fourth argument is float, it must be passed on the stack.

Remaining arguments are passed on the stack. Arguments on the stack should be placed
on the stack in reverse order. Call-by-reference arguments load address (*la* instruction)
and call-by-value load the value.

Calling procedure should use the "jal <proc>" instruction.

Upon completion of the procedure, the caller procedure must restore any saved non-preserved registers and adjust the stack point (**$sp**) as necessary if any arguments were passed on the stack.

Note, for floating-point arguments appearing in registers you must allocate a pair of registers (even if it's a single precision argument) that start with an even register.

8.4 Linkage

The term *linkage* refers to the basic process of getting to a procedure and getting back to the correct location in the calling routine. This does not include argument transmission, which is addressed in the next section.

The basic linkage operation use the **jal** and **jr** instructions. Both instructions utilize the $ra register. This register is set to the return address as part of the procedure call.

The call to a procedure/function requires the procedure/function name, generically labeled as <procName>, as follows:

```
jal     <procName>
```

The **jal**, or jump and link, instruction, will copy the **$pc** into the **$ra** register and jump to the procedure <procName>. Recall that the **$pc** register points to the next instruction to be executed. That will be the instruction immediately after the call, which is the correct place to return to when the procedure/function has completed.

If the procedure/function does not call any other procedures/functions, nothing additional is required with regard to the **$ra** register.

A procedure that does not call another procedure is referred to as a "leaf procedure". A procedure that calls another procedure is referred to as a "non-leaf procedure".

The return from procedure is as follows:

```
jr      $ra
```

If the procedure/function calls yet another procedure/function, the **$ra** must be preserved. Since **$ra** contains the return address, it will be changed when the procedure/function calls the next procedure/function. As such, it must be saved and restored from the stack in the calling procedure. This is typically performed only once at the beginning and then at the end of the procedure (for non-leaf procedures).

Refer to the example programs for a more detailed series of examples that demonstrate the linkage.

8.5 Argument Transmission

Based on the context, parameters may be transmitted to procedures/functions as either values or addresses. These basic approaches are implemented in high-level languages.

The basic argument transmission is accomplished via a combination of registers and the stack.

8.5.1 Call-by-Value

Call-by-value involves passing a copy of the information being passed to the procedure or function. As such, the original value can not be altered.

8.5.2 Call-by-Reference

Call-by-reference involves passing the address of the variables. Call-by-reference is used when passing arrays or when passing variables that will be altered or set by the procedure or function.

8.5.3 Argument Transmission Conventions

The basic argument transmission is accomplished via a combination of registers and the stack.

Integer arguments can be passed in registers **$a0**, **$a1**, **$a2**, and **$a3** and floating-point values passed in **$f12** and **$f14** (single or double precision floating-point).

- The first argument is passed in either **$a0** or **$f12** (**$a0** if integer or **$f12** if float single or double precision).
- The second argument is passed in either **$a1** or **$f14** (**$a1** if integer or **$f14** if float single or double precision).
- The third argument is passed in **$a2** (integer only).
- If the third argument is float, it must be passed on the stack.
- The fourth argument is passed in **$a3** (integer only).
- If the fourth argument is float, it must be passed on the stack.

If the first argument is integer, **$a0** is used and **$f12** should not be used at all. If the first argument is floating-point value, **$f12** is used and **$a0** is not used at all. Any additional arguments are passed on the stack.

The following table shows the argument order and register allocation.

	1st	2nd	3rd	4th	5th	Nth
integer	**$a0**	**$a1**	**$a2**	**$a3**	stack	stack
	or	or	or	or		
floating-point value	**$f12**	**$f14**	stack	stack	stack	stack

Recall that addresses are integers, even when pointing to floating-point values. As such, addresses are passed in integer registers.

8.6 Function Results

A function is expected to return a result (i.e., value returning function).

Integer registers **$v0** or **$v1**/**$v0** are used to return an integer value from a function/procedure call. Floating-point registers **$f0** and **$f1** are used to return a floating-point value from a function/procedure.

8.7 Registers Preservation Conventions

The MIPS calling convention requires that only specific registers (not all) be saved across procedure calls.

- Integer registers **$s0** - **$s7** must be saved by the procedure.
- Floating-point registers **$f20** - **$f30** must be saved by the procedure.

When writing a procedure, this will require that the registers **$s0** - **$s7** or **$f20** - **$f30** (single or double precision) be pushed and popped from the stack if those registers are utilized/changed. When calling a procedure, the main routine must be written so that any values required across procedure calls be placed in register **$s0** - **$s7** or **$f20** - **$f30** (single or double precision).

Integer registers **$t0** - **$t9** and floating-point registers **$f4** - **$f10** and **$f16** - **$f18** (single or double precision) are used to hold temporary quantities that do not need to be preserved across procedure calls.

8.8 Miscellaneous Register Usage

Registers **$at**, **$k0**, and **$k1** are reserved for the assembler and operating system and should not be used by programs. Register **$fp** is used to point to the procedure call frame on the stack. This can be used when arguments are passed on the stack.

Register **$gp** is used as a global point (to point to globally accessible data areas). This register is not typically used when writing assembly programs directly.

8.9 Summary, Callee Conventions

The calling convention addresses specific requirements for the *callee* or routine that is being called from another procedure (which includes the main routine).

- Push any altered "saved" registers on the stack.
  - Specifically, this includes **$s0 - $s7**, **$f20 - $f30**, **$ra**, **$fp**, or **$gp**.
  - If the procedure is a non-leaf procedure, **$ra** must be saved.
  - If **$fp** is altered, **$fp** must be saved which is required when arguments are passed on the stack
  - Space for local variables should be created on the stack for stack dynamic local variables.
- *Note*, when altering the **$sp** register, it should be done in a single operation (instead of a series).
- If arguments are passed on the stack, **$fp** should be set as follows:
  - **$fp = $sp** + (frame size)
  - This will set **$fp** pointing to the first argument passed on the stack.

The procedure can access first 4 integer arguments in registers **$a0 - $a3** and the first two float registers **$f12 - $f14**.

Arguments passed on the stack can be accessed using **$fp**. The procedure should place returned values (if any) into **$v0** and **$v1**.

- Restore saved registers
  - Includes **$s0 - $s7**, **$fp**, **$ra**, **$gp** if they were pushed.
  - Return to the calling procedure via the **jr $ra** instruction.

The procedures example section provides a series of example procedures and functions including register usage and argument transmission.

8.10   Call Frame

The procedure/function call frame or activation record is what the information placed on the stack is called. As noted in the previous sections, the procedure call frame includes passed parameters (if any) and the preserved registers. In addition, space for the procedures' local variables (if any) is allocated on the stack.

A general overview of the call frame is shown as follows:

Call Frame		
		Arguments
		Preserved Registers
		Local Variables

Each part of the call frame may be a different size based on how many arguments are passed (if any), which registers must be preserved (if any), or the amount and size of the local variables (if any).

8.10.1.1   Stack Dynamic Local Variables

The local variables, also referred to as stack dynamic local variables, are typically allocated by the compiler and assigned to stack locations. This allows a more efficient use of memory for high-level languages. This can be very important in large programs.

For example, assume there are 10 procedures each with a locally declared 100,000 element array of integers. Since each integer typically requires 4-bytes, this would mean 400,000 bytes for each procedure with a combined total of 4,000,000 bytes (or about ~4MB) for all ten procedures.

For the standard method of stack dynamic local variables, each array is only allocated when the procedure is active (i.e., being executed). If none of the procedures/functions are called, no memory is allocated. If only two of the arrays are active at any given time, only 800,000 bytes are allocated at any given time.

However, if the arrays were to be declared statically (i.e., not the standard local declaration in the previous examples), the ~4MB of memory is allocated even if none of the procedures are ever called. This can lead to excessive memory usage which can slow a program down.

8.11   Procedure Examples

This section presents a series of example procedures of varying complexity.

8.11.1   Example Program, Power Function

This section presents a very simple example of a function call. The example includes a simple main procedure and a simple function that computes x^y (i.e., x to the y power). The high-level language call, shown in C/C++ here, would be:

```
answer = power(x, y);
```

Where x and y are passed by value and the result is returned to the variable ***answer***. The main passes the arguments by value and receives the result in $v0 (as per the convention). The main then saves the result into the variable ***answer***.

```
#   Example function to demonstrate calling conventions
#   Function computes power (i.e., x to y power).

# ----------------------------------------------------------
#   Data Declarations

.data

x:          .word       3
y:          .word       5
answer:     .word       0

# ------------------------------------
#   Main routine.
#   Call simple procedure to add two numbers.

.text
.globl      main
.ent        main
main:
```

```
        lw    $a0, x                      # pass arg's to function
        lw    $a1, y
        jal   power
        sw    $v0, answer

        li    $v0, 10
        syscall                           # terminate
.end main

# ------------------------------------
# Function to find and return x^y

# -----
#  Arguments
#     $a0 - x
#     $a1 - y
#  Returns
#     $v0 - x^y

.globl      power
.ent        power
power:
      li    $v0, 1
      li    $t0, 0

powLoop:
      mul   $v0, $v0, $a0
      add   $t0, $t0, 1
      blt   $t0, $a1, powLoop

      jr    $ra
.end power
```

Refer to the next section for a more complex example.

8.11.2   Example program, Summation Function

The following is an example program to demonstrate a procedure call.

```
#   Example function to demonstrate calling conventions.
#   Simple function to sum six arguments.
#   ------------------------------------
```

```
# Data Declarations

.data

num1:       .word      3
num2:       .word      5
num3:       .word      3
num4:       .word      5
num5:       .word      3
num6:       .word      5
sum:        .word      0

# ------------------------------------
#  Main routine.
#  Call function to add six numbers.
#     First 4 arguments are passed in $a0-$a3.
#     Next 2 arguments are passed on the stack.

.text

.globl     main
.ent       main
main:
  lw       $a0, num1                # pass arg's
  lw       $a1, num2
  lw       $a2, num3
  lw       $a3, num4
  lw       $t0, num5
  lw       $t1, num6
  subu     $sp, $sp, 8
  sw       $t0, ($sp)
  sw       $t1, 4($sp)
  jal      addem
  sw       $v0, sum

  addu     $sp, $sp, 8              # clear stack

  li       $v0,10
  syscall                          # terminate
.end main

# ------------------------------------
#  Example function to add 6 numbers
```

```
# -----
#  Arguments
#     $a0 - num1
#     $a1 - num2
#     $a2 - num3
#     $a3 - num4
#     ($fp) - num5
#     4($fp) - num6

#  Returns
#     $v0 – num1+num2+num3+num4+num5+num6

.globl     addem
.ent       addem
addem:

  subu     $sp, $sp, 4                 # preserve registers
  sw       $fp, ($sp)

  addu     $fp, $sp, 4                 # set frame pointer

# -----
#  Perform additions.

  li       $v0, 0
  add      $v0, $v0, $a0               # num1
  add      $v0, $v0, $a1               # num2
  add      $v0, $v0, $a2               # num3
  add      $v0, $v0, $a3               # num4
  lw       $t0, ($fp)                  # num5
  add      $v0, $v0, $t0
  lw       $t0, 4($fp)                 # num6
  add      $v0, $v0, $t0

# -----
#  Restore registers.

  lw       $fp, ($sp)
  addu     $sp, $sp, 4

  jr $ra
.end addem
```

Refer to the next section for a more complex example.

8.11.3 Example Program, Pythagorean Theorem Procedure

The following is an example of a procedure that calls another function. Given the *a* and *b* sides of a right triangle, the *c* side can be computed as follows:

$$cSide = \sqrt{aSide^2 + bSide^2}$$

b

a

This example program will call a procedure to compute the *c* sides of a series of right triangles. The *a* sides and *b* sides are stored in an arrays, *aSides[]* and *bSides[]* and results stored into an array, *cSides[]*. The procedure will also compute the minimum, maximum, sum, and average of the *cSides[]* values. All values are integers. In order to compute the integer square root, a *iSqrt()* function is used. The *iSqrt()* function uses a simplified version of Newton's method.

```
#   Example program to calculate the cSide for each
#   right triangle in a series of right triangles
#   given the aSides and bSides using the
#   Pythagorean theorem.

#   Pythagorean theorem:
#      cSide = sqrt ( aSide^2 + bSide^2 )

#   Provides examples of MIPS procedure calling.

# --------------------------------------------------------
#   Data Declarations

.data

aSides:     .word   19, 17, 15, 13, 11, 19, 17, 15, 13, 11
            .word   12, 14, 16, 18, 10
bSides:     .word   34, 32, 31, 35, 34, 33, 32, 37, 38, 39
            .word   32, 30, 36, 38, 30
cSides:     .space  60

length:     .word   15

min:        .word   0
max:        .word   0
sum:        .word   0
ave:        .word   0
```

```
# ------------------------------------------------------
#  text/code section

.text
.globl     main
.ent       main
main:

# -----
#  Main program calls the cSidesStats routine.
#  The HLL call is as follows:
#     cSidesStats(aSides, bSides, cSides, length, min,
#                      max, sum, ave)
#  Note:
#     The arrays are passed by reference
#     The length is passed by value
#     The min, max, sum, and ave are pass by reference.

     la    $a0, aSides               # address of array
     la    $a1, bSides               # address of array
     la    $a2, cSides               # address of array
     lw    $a3, length               # value of length

     la    $t0, min                  # address for min
     la    $t1, max                  # address for max
     la    $t2, sum                  # address for sum
     la    $t3, ave                  # address for ave

     subu $sp, $sp, 16
     sw    $t0, ($sp)                # push addresses
     sw    $t1, 4($sp)
     sw    $t2, 8($sp)
     sw    $t3, 12($sp)

     jal  cSidesStats                # call routine
     addu $sp, $sp, 16               # clear arguments

# -----
#  Done, terminate program.

     li    $v0, 10                   # terminate
     syscall                         # system call
.end main
```

```
# -------------------------------------------------------
#   Function to calculate the cSides[] for each right
#   triangle in a series of right triangles given the
#   aSides[] and bSides[] using the Pythagorean theorem.

#   Pythagorean theorem formula:
#       cSides[n] = sqrt ( aSides[n]^2 + bSides[n]^2 )

#   Also finds and returns the minimum, maximum, sum,
#   and average for the cSides.
#   Uses the iSqrt() routine to find the integer
#   square root of an integer.

# -----
#   Arguments:
#       $a0 - address of aSides[]
#       $a1 - address of bSides[]
#       $a2 - address of cSides[]
#       $a3 - list length
#       ($fp) - addr of min
#       4($fp) - addr of max
#       8($fp) - addr of sum
#       12($fp) - addr of ave

#   Returns (via passed addresses):
#       cSides[]
#       min
#       max
#       sum
#       ave

.globl      cSidesStats
.ent        cSidesStats
cSidesStats:
    subu $sp, $sp, 24               # preserve registers
    sw   $s0, 0($sp)
    sw   $s1, 4($sp)
    sw   $s2, 8($sp)
    sw   $s3, 12($sp)
    sw   $fp, 16($sp)
    sw   $ra, 20($sp)

    addu $fp, $sp, 24               # set frame pointer
```

```
# -----
#  Loop to calculate cSides[]
#     Note, must use $s<n> registers due to iSqrt() call

      move  $s0, $a0                # address of aSides
      move  $s1, $a1                # address of bSides
      move  $s2, $a2                # address of cSides
      li    $s3, 0                  # index = 0

cSidesLoop:
      lw    $t0, ($s0)              # get aSides[n]
      mul   $t0, $t0, $t0           # aSides[n]^2
      lw    $t1, ($s1)              # get bSides[n]
      mul   $t1, $t1, $t1           # bSides[n]^2
      add   $a0, $t0, $t1
      jal   iSqrt                   # call iSqrt()

      sw    $v0, ($s2)              # save to cSides[n]

      addu  $s0, $s0, 4             # update aSides addr
      addu  $s1, $s1, 4             # update bSides addr
      addu  $s2, $s2, 4             # update cSides addr
      addu  $s3, $s3, 1             # index++

      blt   $s3, $a3, cSidesLoop    # if indx<len, loop

# -----
#  Loop to find minimum, maximum, and sum.

      move  $s2, $a2                # strt addr of cSides
      li    $t0, 0                  # index = 0
      lw    $t1, ($s2)              # min = cSides[0]
      lw    $t2, ($s2)              # max = cSides[0]
      li    $t3, 0                  # sum = 0

statsLoop:
      lw    $t4, ($s2)              # get cSides[n]

      bge   $t4, $t1, notNewMin     # if cSides[n] >=
                                    #   item -> skip
      move  $t1, $t4                # set new min value
notNewMin:
```

```
        ble   $t4, $t2, notNewMax         # if cSides[n] <=
                                          #   item -> skip
        move  $t2, $t4                    # set new max value
notNewMax:

        add   $t3, $t3, $t4               # sum += cSides[n]

        addu  $s2, $s2, 4                 # update cSides addr
        addu  $t0, $t0, 1                 # index++

        blt   $t0, $a3, statsLoop         # if indx<len, loop

        lw    $t5, ($fp)                  # get address of min
        sw    $t1, ($t5)                  # save min

        lw    $t5, 4($fp)                 # get address of max
        sw    $t2, ($t5)                  # save max

        lw    $t5, 8($fp)                 # get address of sum
        sw    $t3, ($t5)                  # save sum

        div   $t0, $t3, $a3               # ave = sum / len

        lw    $t5, 12($fp)                # get address of ave
        sw    $t0, ($t5)                  # save ave

    # -----
    #  Done, restore registers and return to calling routine.

        lw    $s0, 0($sp)
        lw    $s1, 4($sp)
        lw    $s2, 8($sp)
        lw    $s3, 12($sp)
        lw    $fp, 16($sp)
        lw    $ra, 20($sp)
        addu  $sp, $sp, 24
        jr    $ra
.end cSidesStats

    # -------------------------------------------------------
    #  Function to compute integer square root for
    #  an integer value.
```

```
#   Uses a simplified version of Newtons method.
#      x = N
#      iterate 20 times:
#           x' = (x + N/x) / 2
#           x = x'

# -----
#   Arguments
#      $a0 - N
#   Returns
#      $v0 - integer square root of N

.globl    iSqrt
.ent      iSqrt
iSqrt:
     move $v0, $a0                  # $v0 = x = N
     li   $t0, 0                    # counter
sqrLoop:
     div  $t1, $a0, $v0             # N/x
     add  $v0, $t1, $v0             # x + N/x
     div  $v0, $v0, 2               # (x + N/x)/2

     add  $t0, $t0, 1
     blt  $t0, 20, sqrLoop

     jr   $ra
.end iSqrt
```

This example uses a simplified version of Newton's method. Further improvements are
left to the reader as an exercise.

9.0   QtSpim System Service Calls

The operating system must provide some basic services for functions that a user program can not easily perform on its own. Some key examples include input and output operations. These functions are typically referred to as *system services*. The QtSpim simulator provides a series of operating system like services by using a **syscall** instruction.

To request a specific service from the QtSpim simulator, the 'call code' is loaded in the **$v0** register. Based on the specific system service being requested, additional information may be needed which is loaded in the argument registers (as noted in the Procedures/Functions section).

9.1 Supported QtSpim System Services

A list of the supported system services is listed in the below table. A series of examples are provided in the following sections.

Service Name	Call Code	Input	Output
Print Integer (32-bit)	1	**$a0** → integer to be printed	
Print Float (32-bit)	2	**$f12** → 32-bit floating-point value to be printed	
Print Double (64-bit)	3	**$f12** → 64-bit floating-point value to be printed	
Print String	4	**$a0** → starting address of NULL terminated string to be printed	
Read Integer (32-bit)	5		**$v0** → 32-bit integer entered by user
Read Float (32-bit)	6		**$f0** → 32-bit floating-point value entered by user

Read Double (64-bit)	7		$f0 → 64-bit floating-point value entered by user
Read String	8	$a0 → starting address of buffer (of where to store character entered by user) $a1 → length of buffer	
Allocate Memory	9	$a0 → number of bytes to allocate	$v0 → starting address of allocated memory
Terminate	10		
Print Character	11	$a0 → character to be printed	
Read Character	12		$v0 → character entered by user
File Open	13	$a0 → file name string, NULL terminated $a1 → access flags $a2 → file mode, (UNIX style)	$v0 → file descriptor
File Read	14	$a0 → file descriptor $a1 → buffer starting address $a2 → number of bytes to read	$v0 → number of bytes actually read from file (-1 = error, 0 = end of file)
File Write	15	$a0 → file descriptor $a1 → buffer starting address $a2 → number of bytes to read	$v0 → number of bytes actually written to file (-1 = error, 0 = end of file)
File Close	16	$a0 → file descriptor	

The file open access flags are defined as follows:

```
Read = 0x0, Write = 0x1, Read/Write = 0x2
OR Create = 0x100, Truncate = 0x200, Append = 0x8
OR Text = 0x4000, Binary = 0x8000
```

For example, for a file read operation, the 0x0 would be selected. For a file write operation, the 0x1 would be selected.

9.2 QtSpim System Services Examples

This section provides a series of examples using system service calls.

The system service calls follow the standard calling convention in that the temporary registers (**$t0** - **$t9**) may be altered and the saved registers (**$s0** - **$s7**, **$fp**, **$ra**) will be preserved. As such, if a series of values is being printed in a loop, a saved register would be required for the loop counter and the current array address/index.

9.2.1 Example Program, Display String and Integer

The following code provides an example of how to display a string and an integer.

```
#  Example program to display a string and an integer.
#  Demonstrates use of QtSpim system service calls.
#  -------------------------------------------------------
#  Data Declarations

.data
hdr:          .ascii     "Example\n"
              .asciiz    "The meaning of life is: "
number:       .word      42

#  -------------------------------------------------------
#  text/code section

.text
.globl     main
.ent       main
main:
    la    $a0, hdr              # addr of NULL
                                #   terminated string
    li    $v0, 4                # call code, print string
    syscall                     # system call

    li    $v0, 1                # call code, print int
    lw    $a0, number           # value for int to print
    syscall                     # system call

# -----
#  Done, terminate program.
```

```
        li    $v0, 10         # terminate
        syscall               # system call
.end main
```

Note, in this example, the string definition ensures the NULL termination as required by the system service.

The output for the example would be displayed to the QtSpim console window. For example:

The console window can be displayed or hidden from the Windows menu (on the top bar).

9.2.2 Example Program, Display Array

This section provides an example of how to display an array. In this example, an array of numbers is displayed to the screen with five numbers per line (arbitrarily chosen) to make the output appear more pleasing.

Since the system service call is utilized for the print function, the saved register must be used. Refer to the Procedures/Functions section for additional information regarding the MIPS calling conventions.

```
    #  Example program to display an array.
    #  Demonstrates use of QtSpim system service calls.


    # -------------------------------------------------------
    #  Data Declarations

    .data
    hdr:        .ascii      "Array Values\n"
                .asciiz     "-----------------------\n\n"
    spaces:     .asciiz     "   "
    newLine:    .asciiz     "\n"
```

```
array:      .word       11, 13, 15, 17, 19
            .word       21, 23, 25, 27, 29
            .word       31, 33, 35, 37, 39
            .word       41, 43, 45, 47
length:     .word       19

# -------------------------------------------------------
#  text/code section

.text
.globl      main
.ent        main
main:
    li    $v0, 4               # print header string
    la    $a0, hdr
    syscall

    la    $s0, array
    li    $s1, 0
    lw    $s2, length

printLoop:
    li    $v0, 1               # call code for print int
    lw    $a0, ($s0)           # get array[i]
    syscall                    # system call

    li    $v0, 4               # print spaces
    la    $a0, spaces
    syscall

    addu $s0, $s0, 4           # update addr (next word)
    add  $s1, $s1, 1           # increment counter

    rem  $t0, $s1, 5
    bnez $t0, skipNewLine

    li    $v0, 4               # print new line
    la    $a0, newLine
    syscall

skipNewLine:
    bne  $s1, $s2, printLoop   # if cnter<len -> loop
```

```
        # -----
        #  Done, terminate program.

            li    $v0, 10                  # terminate
            syscall                        # system call
        .end main
```

The output for the example would be displayed to the QtSpim console window.

For example:

This example program does not align the values (when printed). The values only appear aligned since they all have the same number of digits.

9.2.3 Example Program, Read Integer

This section provides an example of how to display a prompt string, read an integer value, square that integer value, and display the final result.

It must be noted that the QtSpim read integer system service is fairly basic and does not perform error checking or handle backspace/delete. As such, the number must be entered correctly by the user. If invalid numbers, such as (a12 or 12q34) are entered, the input will be mis-interpreted resulting in unexpected or invalid values.

If desired, the numeric input can be read as a string and converted into an integer with the appropriate error handling. This is left to the user as an exercise.

```
        #  Example program to display an array.
        #  Demonstrates use of QtSpim system service calls.

        # -------------------------------------------------------
        #  Data Declarations
        .data
```

```
hdr:          .ascii       "Squaring Example\n"
              .asciiz      "Enter Value: "
ansMsg:       .asciiz      "Value Squared: "
value:        .word        0

# --------------------------------------------------------
#  text/code section

.text
.globl     main
.ent       main
main:
     li    $v0, 4          # call code for print string
     la    $a0, hdr        # addr of NULL terminated str
     syscall               # system call

     li    $v0, 5          # call code for read integer
     syscall               # system call (result in $v0)

     mul   $t0, $v0, $v0   # square answer
     sw    $t0, value      # save to variable

     li    $v0, 4          # call code for print string
     la    $a0, ansMsg     # addr of NULL terminated str
     syscall               # system call

     li    $v0, 1          # call code for print integer
     lw    $a0, value      # value for integer to print
     syscall               # system call

# -----
#  Done, terminate program.

     li    $v0, 10         # terminate
     syscall               # system call
.end main
```

The output for the example would be displayed to the QtSpim console window. For example:

The console window must be selected in order to enter input. *Note*, the default console window size will typically be larger than what is shown above.

9.2.4 Example Program, Read String

This section provides an example of how to display a prompt string and read a string of characters. As previously noted, at the assembly level, strings are a series of contiguously defined byte-sized characters, typically terminated with a NULL byte (0x00).

In order to read a string, some space for where to place the characters read must be created. The QtSpim system service for read string will always terminate the string with a NULL byte which must be accommodated for in the space allocated.

In this example, a variable, *userAns*, was defined with fifty-two (52) bytes of space. This allows up to fifty (50) characters, a line feed (0x0A), and the NULL termination. It should be noted that if fifty-one (51) characters are entered, the input will be automatically terminated, *without the user pressing enter*, and the NULL added to the string. This can very awkward when users are entering input, so input string sizes should be chosen carefully.

When the QtSpim system service is called, the string address (in **$a0**) and length (in **$a1**) must be provided. It is important that the correct length be provided as an error could result in memory, and thus other variables, being over-written. Such problems can be very difficult to find as the symptom will typically be in a different location than the actual problem.

```
# Example program to demonstrate string input
# ----------------------------------------------------
#  Data Declarations
```

```
          .data
hdr:          .ascii      "Reading Characters Example\n\n"
              .asciiz     "Enter Your Name: "
hiMsg:        .asciiz     "\nHello, "
userAns:      .space      52

# ------------------------------------------------------
#  text/code section

.text
.globl      main
.ent        main
main:
      li    $v0, 4               # call code, print string
      la    $a0, hdr             # addr of string
      syscall                    # system call

      li    $v0, 8               # call code, read string
      la    $a0, userAns         # addr, where to put chars
      li    $a1, 52              # max chars for string
      syscall                    # system call

      li    $v0, 4               # call code, print string
      la    $a0, hiMsg           # address string
      syscall                    # system call

      li    $v0, 4               # call code, print string
      la    $a0, userAns             # address string
      syscall                    # system call

      li    $v0, 10              # call code for terminate
      syscall                    # system call
.end main
```

The output and input for the example would be displayed to the QtSpim console
window.

For example:

The console window must be selected in order to enter input. *Note*, the default console window size will typically be larger than what is shown above.

10.0   Multi-dimension Array Implementation

This chapter provides a summary of the implementation of multiple dimension arrays as viewed from assembly language.

Memory is inherently a single dimension entity. As such, a multi-dimension array is implemented as sets of single dimension array. There are two primary ways this can be performed; row-major and column-major. Each is explained in subsequent sections.

To simplify the explanation, this section focuses on two-dimensional arrays. The general process extends to higher dimensions.

10.1   High-Level Language View

Multi-Dimension arrays are sometimes used in high-level languages. For example, in C/C++, the declaration of: `int arr [3][4]` would declare an array as follows:

	arr[2][0]	arr[2][1]	arr[2][2]	arr[2][3]
	arr[1][0]	arr[1][1]	arr[1][2]	arr[1][3]
arr	arr[0][0]	arr[0][1]	arr[0][2]	arr[0][3]

It is expected that the reader is generally familiar with the high-level language use of two-dimensional arrays.

10.2   Row-Major

Row-major assigns each row as a single dimension array in memory, one row after the next until all rows are in memory.

The formula to convert two-dimensional array indexes (row, column) into a single dimension, row-major memory offset is as follows:

```
addr = baseAddr + (rowIdx * numOfCols + colIdx) * dataSize
```

Where the *base address* is the starting address of the array, *dataSize* is the size of the data in bytes, and *numOfCols* is the dimension or number of the columns in the two-dimension array. In this example, the number of columns in the array is 4 (from the previous high-level language declaration).

For example, to access the **arr[1][2]** element (labeled '6' in the above diagram), assuming the array is composed of 32-bit sized elements it would be:

```
address  = arr + (1 * 4 + 2) * 4 = arr + (4 + 2) * 4
         = arr + 6 * 4 = arr + 24
```

Which generates the correct, final address.

10.3   Column-Major

Column-major assigns each column as a single dimension array in memory, one column after the next until all rows are in memory.

11	arr[2][3]
10	arr[1][3]
9	arr[0][3]
8	arr[2][2]
7	arr[1][2]
6	arr[0][2]
5	arr[2][1]
4	arr[1][1]
3	arr[0][1]
2	arr[2][0]
1	arr[1][0]
0	arr[0][0]

arr

2	5	8	11
1	4	7	10
0	3	6	9

The formula to convert two-dimensional array indexes (row, column) into a single dimension, column-major memory offset is as follows:

```
addr = baseAddr + (colIdx * numOfRows + rowIdx) * dataSize
```

Where the base *address* is the starting address of the array, *dataSize* is the size of the data in bytes, and *numOfRows* is the dimension or number of the rows in the two-dimension array. In this example, the number of rows in the array is 3 (from the previous high-level language declaration).

For example, to access the `arr[1][2]` element (labeled '7' in the above diagram), assuming the array is composed of 32-bit sized elements it would be:

```
address  = arr + (2 * 3 + 1) * 4 = arr + (6 + 1) * 4
         = arr + 7 * 4 = arr + 28
```

Which generates the correct, final address.

10.4   Example Program, Matrix Diagonal Summation

The following code provides an example of how to access elements in a two-dimensional array. This example adds the elements on the diagonal of a two-dimensional array.

For example, given the logical view of a five-by-five square matrix:

11	12	13	14	15
16	17	18	19	20
21	22	23	24	25
26	27	28	29	30
31	32	33	34	35

The main diagonal contains the numbers, 11, 17, 23, 29, and 35.

```
#   Example program to compute the sum of diagonal
#   in a square two-dimensional, row-major array
#   Demonstrates multi-dimension array indexing.
#   Assumes row-major ordering.
#   -------------------------------------------------------
#   Data Declarations

.data
mdArray:    .word       11, 12, 13, 14, 15
            .word       16, 17, 18, 19, 20
            .word       21, 22, 23, 24, 25
            .word       26, 27, 28, 29, 30
            .word       31, 32, 33, 34, 35

size:       .word       5
dSum:       .word       0

DATASIZE = 4                            # 4 bytes for words

finalMsg:   .ascii      "Two-Dimensional Diagonal"
            .ascii      "Summation\n\n"
            .asciiz     "Diagonal Sum = "

#   -------------------------------------------------------
```

```
#   Text/code section

.text
.globl      main
.ent        main
main:

# -----
#  Call function to sum the diagonal
#     (of square two-dimensional array)

      la    $a0, mdArray          # base address of array
      lw    $a1, size             # array size

      jal   diagSummer
      sw    $v0, dSum

# -----
#  Display final result.

      li    $v0, 4                # print prompt string
      la    $a0, finalMsg
      syscall

      li    $v0, 1                # print integer
      lw    $a0, dSum
      syscall

# -----
# Done, terminate program.

      li    $v0, 10               # terminate
      syscall                     # system call
.end main

# ------------------------------------------------------
#  Simple function to sum the diagonals of a
#  square two-dimensional array.

#  Approach
#    loop i = 0 to len-1
#         sum = sum + mdArray[i][i]
#  Note, for two-dimensional array:
```

```
#    addr = baseAddr + (rowIdx * numOfCols + colIdx)
#                                     * dataSize
#  Since the two-dimensional array is given as square,
#  the row and column dimensions are the same size.
#  -----
#  Arguments
#    $a0 - array base address
#    $a1 - size (of square two-dimension array)
#  Returns
#    $v0 - sum of diagonals

.globl      diagSummer
.ent        diagSummer
diagSummer:

        li   $v0, 0                 # sum=0
        li   $t1, 0                 # loop index, i=0

diagSumLoop:
        mul  $t3, $t1, $a1          # (rowIdx * colSize
        add  $t3, $t3, $t1          #            + colIdx)
                                    # note, rowIdx=colIdx
        mul  $t3, $t3, DATASIZE     #            * dataSize
        add  $t4, $a0, $t3          # + base address

        lw   $t5, ($t4)             # get mdArray[i][i]

        add  $v0, $v0, $t5          # sum = sum+mdArray[i][i]


        add  $t1, $t1, 1            # i = i + 1
        blt  $t1, $a1, diagSumLoop

# -----
#  Done, return to calling routine.

        jr   $ra
.end diagSummer
```

While not mathematically useful, this does demonstrate how elements in a two-dimensional array are accessed.

11.0   Recursion

The Google search result for recursion, shows "*Did you mean: Recursion*".

Recursion is the idea that a function may call itself (which is the basis for the joke). Recursion is a powerful general-purpose programming technique and is used for some important applications including searching and sorting.

Recursion can be very confusing in its simplicity. The simple examples in this section will not be enough in themselves for the reader to obtain recursive enlightenment. The goal of this section is to provide some insight into the underlying mechanisms that support recursion. The simple examples here which are used to introduce recursion are meant to help demonstrate the form and structure for recursion. More complex examples (than will be discussed here) should be studied and implemented in order to ensure a complete appreciation for the power of recursion.

The procedure/function calling process previously described supports recursion without any changes.

A recursive function must have a recursive definition that includes:

1.   Base case, or cases, that provide a simple result (that defines when the recursion should stop).

2.   Rule, or set of rules, that reduce toward the base case.

This definition is referred to as a recursive relation.

11.1   Recursion Example, Factorial

The factorial function is mathematically defined as follows:

$$n! \ = \ \prod_{k=1}^{n} k$$

Or more familiarly, you might see **5!** as:

$$5! \ = \ 5 \times 4 \times 3 \times 2 \times 1$$

It must be noted that this function could easily be computed with a loop. However, the reason this is done recursively is to provide a simple example of how recursion works.

Chapter 11.0 ◄ Recursion

A typical recursive definition for factorial is:

$$factorial(n) = \begin{cases} 1 & \text{if } n = 0 \\ n \times factorial(n-1) & \text{if } n \geq 1 \end{cases}$$

This definition assumes that the value of *n* is positive.

11.1.1   Example Program, Recursive Factorial Function

The following code provides an example of the recursive factorial function.

```
#  Example program to demonstrate recursion.
# ---------------------------------------------------------
#  Data Declarations

.data

prompt:     .ascii      "Factorial Example Program\n\n"
            .asciiz     "Enter N value: "
results:    .asciiz     "\nFactorial of N = "

n:          .word 0
answer:     .word 0

# ---------------------------------------------------------
#  Text/code section

.text
.globl      main
.ent        main
main:

# -----
#  Read n value from user

    li    $v0, 4                # print prompt string
    la    $a0, prompt
    syscall

    li    $v0, 5                # read N (as integer)
```

```
        syscall
        sw    $v0, n

#  -----
#  Call factorial function.

        lw    $a0, n
        jal   fact
        sw    $v0, answer

#  -----
#  Display result

        li    $v0, 4                # print prompt string
        la    $a0, results
        syscall

        li    $v0, 1                # print integer
        lw    $a0, answer
        syscall

#  -----
#  Done, terminate program.

        li    $v0, 10               # call code for terminate
        syscall                     # system call
.end main

#  -------------------------------------------------------
#  Factorial function
#  Recursive definition:
#     = 1                if n = 0
#     = n * fact(n-1)    if n >= 1

#  -----
#  Arguments
#     $a0 - n
#  Returns
#     $v0 set to n!

.globl    fact
.ent      fact
fact:
```

```
        subu $sp, $sp, 8
        sw    $ra, ($sp)
        sw    $s0, 4($sp)

        li    $v0, 1                    # check base case
        beq   $a0, 0, factDone

        move  $s0, $a0                  # fact(n-1)
        sub   $a0, $a0, 1
        jal   fact

        mul   $v0, $s0, $v0             # n * fact(n-1)

   factDone:
        lw    $ra, ($sp)
        lw    $s0, 4($sp)
        addu  $sp, $sp, 8
        jr    $ra
   .end  fact
```

The output for the sample program would be displayed to the QtSpim console window.
For example:

Refer to the next section for an explanation of how this function works.

11.1.2   Recursive Factorial Function Call Tree

In order to help understand recursion, a recursion tree can show how the recursive calls interact.

When the initial call occurs from main, the main will start into the *fact()* function (shown as step 1). Since the argument of 5 is not a base case, the *fact()* function must call *fact()* again with the argument of *n*-1 or 4 in this example (step 2). And, again, since 4 is not the base case, the *fact()* function must call *fact()* again with the argument of *n*-1 or 3 in this example (step 3).

This process continues until the argument passed into the *fact()* function meets the base case which is when the argument is equal to 1 (shown as step 5). When this occurs, only then is a return value provided to the previous call (step 6). This return argument is then

used to calculate the previous multiplication which is 2 times 1 which will return a value to the previous call (as shown in step 7).

These returns will continue (steps 8, 9, and 10) until the main has a final answer.

Since the code being executed is the same, each instance of the *fact()* function is different from any other instance only in the arguments and temporary values. The arguments and temporary values for each instance are different since they are maintained on the stack as required by the standard calling convention.

For example, consider a call to factorial with **n** = 2 (step 4 on the diagram). The return address, **$ra**, and previous contents of **$s0** are preserved by pushing them on the stack in accordance with the standard calling convention. The base case is checked and since **n** ≠ 1 it continues to save the original value of 1 into **$s0**, decrements the original argument, **n,** by 1 and calls the *fact()* function (with **n** = 1). The call for the *fact()* function (step 5 in the diagram) is like any other function call in that it must follow the standard calling convention, which requires preserving **$ra** and **$s0** (since they are changed). This is when the function returns an answer, 1 in this specific case, that answer in **$v0** is then multiplied by the original **n** value in **$s0** and returned to the calling routine.

As such, the foundation for recursion is the procedure call frame or activation record. In general, it can be simply stated that recursion is stack-based.

It should also be noted that the height of the recursion tree is directly associated with the amount of stack memory used by the function.

11.2   Recursion Example, Fibonacci

The Fibonacci function is mathematically defined as follows:

$$F_n \; = \; F_{n-1} + F_{n-2}$$

for positive integers with seed values of $F_0 = 0$ and $F_1 = 1$ by definition.

As such, starting from 0 the first 14 numbers in the Fibonacci series are:

0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233

It must be noted that this function could easily be computed with a loop. However, the reason this is done recursively is to provide a simple example of how recursion works.

For example, a typical recursive definition for Fibonacci is:

$$fib(n) = \begin{cases} 0 & \textit{if } n=0 \\ 1 & \textit{if } n=1 \\ fib(n-1) + fib(n-2) & \textit{if } n>1 \end{cases}$$

This definition assumes that the value of **n** is positive.

11.2.1   Example Program, Recursive Fibonacci Function

The following code provides an example of the recursive Fibonacci function.

```
#  Recursive Fibonacci program to demonstrate recursion.
# --------------------------------------------------------
#  Data Declarations

.data
prompt:     .ascii      "Fibonacci Example Program\n\n"
            .asciiz     "Enter N value: "

results:    .asciiz     "\nFibonacci of N = "

n:          .word 0
answer:     .word 0

# --------------------------------------------------------
#  Text/code section

.text
.globl      main
.ent        main
main:

# -----
#  Read n value from user

    li    $v0, 4                  # print prompt string
    la    $a0, prompt
    syscall
```

```
        li    $v0, 5                # read N (as integer)
        syscall
        sw    $v0, n

# -----
#  Call Fibonacci function.

        lw    $a0, n
        jal   fib
        sw    $v0, answer

# -----
#  Display result

        li    $v0, 4                # print prompt string
        la    $a0, results
        syscall

        li    $v0, 1                # print integer
        lw    $a0, answer
        syscall

# -----
#  Done, terminate program.
        li    $v0, 10               # terminate
        syscall                     # system call
.end main

# -------------------------------------------------------
#  Fibonacci function
#  Recursive definition:
#     = 0                     if n = 0
#     = 1                     if n = 1
#     = fib(n-1) + fib(n-2)   if n > 2

# -----
#  Arguments
#     $a0 - n
#  Returns
#     $v0 set to fib(n)

.globl     fib
.ent       fib
```

```
fib:
      subu $sp, $sp, 8
      sw   $ra, ($sp)
      sw   $s0, 4($sp)

      move $v0, $a0              # check for base cases
      ble  $a0, 1, fibDone

      move $s0, $a0              # get fib(n-1)
      sub  $a0, $a0, 1
      jal  fib

      move $a0, $s0
      sub  $a0, $a0, 2          # set n-2
      move $s0, $v0             # save fib(n-1)
      jal  fib                 # get fib(n-2)

      add  $v0, $s0, $v0       # fib(n-1)+fib(n-2)

   fibDone:
      lw   $ra, ($sp)
      lw   $s0, 4($sp)
      addu $sp, $sp, 8
      jr   $ra
.end fib
```
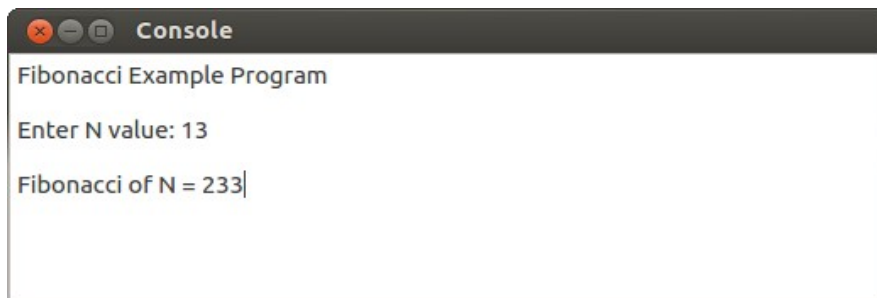
The output for the example would be displayed to the QtSpim console window.

For example:

Refer to the next section for an explanation of how this function works.

11.2.2   Recursive Fibonacci Function Call Tree

The Fibonacci recursion tree appears more complex than the previous factorial tree since the Fibonacci function uses two recursive calls. However, the general process and use of the stack for arguments and temporary values is the same.

As noted in the factorial example, the basis of recursion is the stack. In this example, since two recursive calls are made, the first call will make another call, which may make yet another call. In this manner, the call sequence will follow the order shown in the following diagram.

The following is an example of the call tree for a Fibonacci call with *n* = 4.
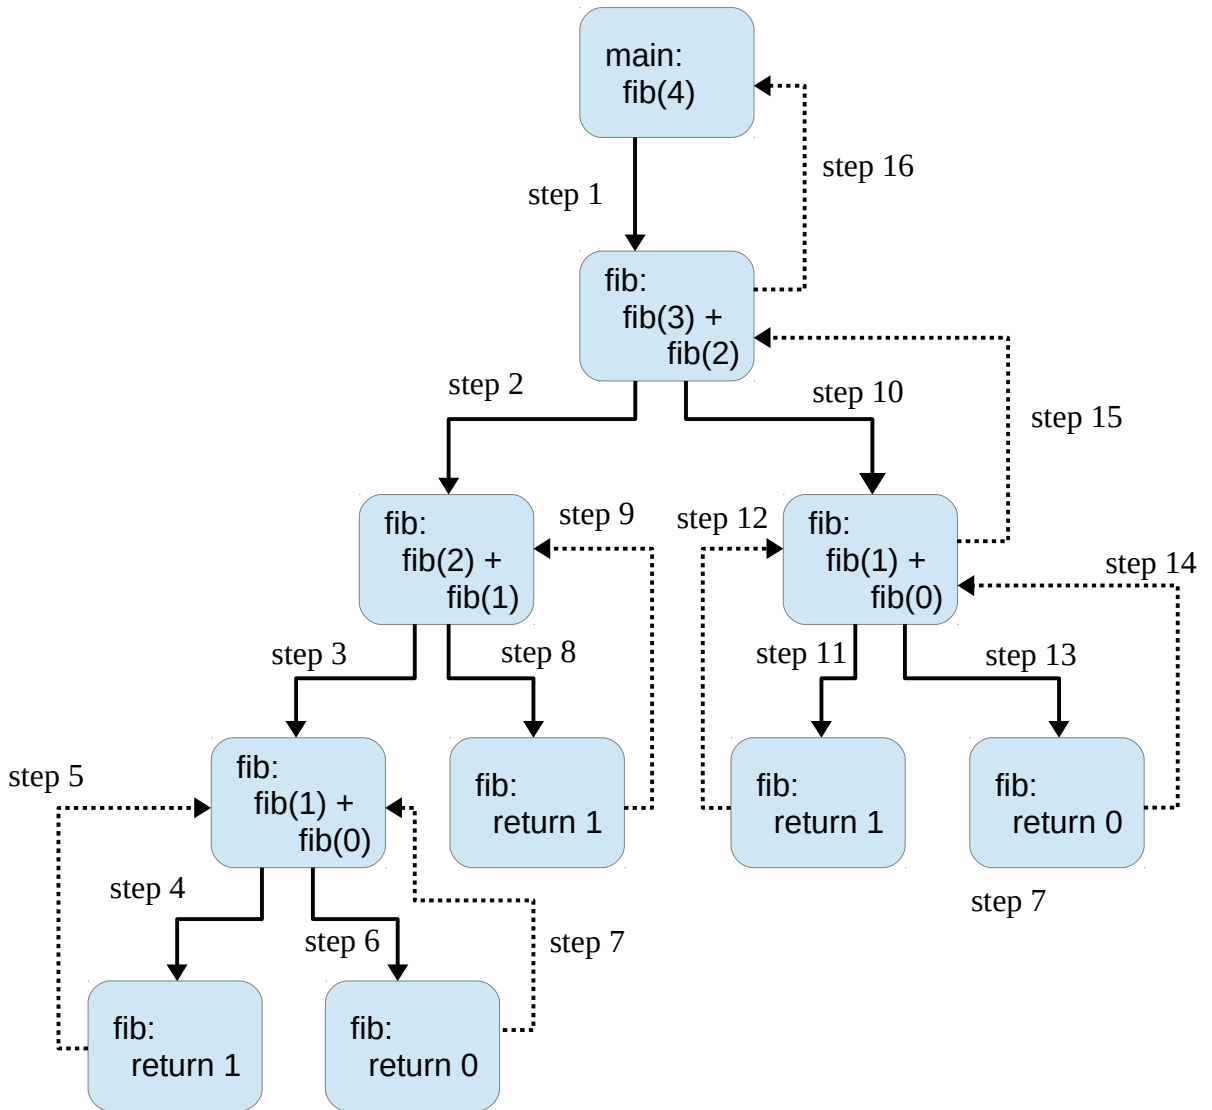
The calls are shown with a solid line and the returns are shown with a dashed line.

Chapter 11.0 ◄ Recursion

12.0   Appendix A – Example Program

Below is a simple example program. This program can be used to test the simulator installation and as an example of the required program formatting.

```
# Example program to find the minimum and maximum from
# a list of numbers.

# -----------------------------------------------------
#  data segment

.data

array:      .word       13, 34, 16, 61, 28
            .word       24, 58, 11, 26, 41
            .word       19,  7, 38, 12, 13
len:        .word       15

hdr:        .ascii      "\nExample program to find max and"
            .asciiz     " min\n\n"
newLine:    .asciiz     "\n"
a1Msg:      .asciiz     "min = "
a2Msg:      .asciiz     "max = "

# -----------------------------------------------------
#  text/code segment
#  QtSpim requires the main procedure be named "main".

.text
.globl      main
.ent        main
main:

#  This program will use pointers.
#     t0 - array address
#     t1 - count of elements
#     s2 - min
#     s3 - max
```

Appendix A – Example Program

```
#     t4 - each word from array
# -----
#  Display header
#  Uses print string system call

        la    $a0, hdr
        li    $v0, 4
        syscall                      # print header

# -----
#  Find max and min of the array.
#   Set min and max to first item in list and then
#   loop through the array and check min and max
#   against each item in the list, updating the min
#   and max values as needed.

        la    $t0, array             # $t0 addr of array
        lw    $t1, len               # $t1 to length
        lw    $s2, ($t0)             # min, $s2 to array[0]
        lw    $s3, ($t0)             # max, $s3 to array[0]

loop:
        lw    $t4, ($t0)             # get array[n]

        bge   $t4, $s2, NotMin# is new min?
        move  $s2, $t4               # set new min

NotMin:
        ble   $t4, $s3, NotMax# is new max?
        move  $s3, $t4               # set new max
NotMax:
        sub   $t1, $t1, 1            # decrement counter
        addu  $t0, $t0, 4            # increment addr by word
        bnez  $t1, loop

# -----
#  Display results min and max.
#   First display string, then value, then a print a
#   new line (for formatting). Do for each max and min.

        la    $a0, a1Msg
        li    $v0, 4
        syscall                      # print "min = "
```

```
        move $a0, $s2
        li   $v0, 1
        syscall                    # print min

        la   $a0, newLine          # print a newline
        li   $v0, 4
        syscall

        la   $a0, a2Msg
        li   $v0, 4
        syscall                    # print "max = "

        move $a0, $s3
        li   $v0, 1
        syscall                    # print max

        la   $a0, newLine          # print a newline
        li   $v0, 4
        syscall

# -----
#  Done, terminate program.
        li   $v0, 10
        syscall                    # all done!

        .end main
```

Appendix A – Example Program

13.0   Appendix B – QtSpim Tutorial

This QtSpim Tutorial is designed to prepare you to use the QtSpim simulator and complete your MIPS assignments more easily.

13.1   Downloading and Installing QtSpim

The first step is to download and install QtSpim for your specific machine. QtSpim is available for Windows, Linux, and MAC OS's.

13.1.1   QtSpim Download URLs

The following are the current URLs for QtSpim.

The QtSpim home page is located at:

> http://spimsimulator.sourceforge.net/

The specific download site is located at:
> http://sourceforge.net/projects/spimsimulator/files/

At the download site there are multiple versions for different target machines. These include Windows (all versions), Linux/Ubuntu (32-bit), Linux/Ubuntu (64-bit), and Mac OS (all versions). Download the latest version for your machine.

These URLs are subject to change. If they do not work, a Google search will find the correct URLs.

13.1.2   Installing QtSpim

Once the package is downloaded, follow the standard installation process for the specific OS being used. This typically will involve double-clicking the downloaded installation package and following the instructions. You will need administrator privileges to perform the installation. Additionally, some installations will require Internet access during the installation.

13.2　　　Working Directory

Create a working directory for the QtSpim assembly source files. This directory can be named anything, but must be legal on the chosen operating system.

13.3　　　Sample Program

Copy or type the provided example program (from Appendix A) to a file in your working directory. This file will be used in the remainder of the tutorial. It demonstrates assembler directives, procedure calls, console I/O, program termination, and good programming practice. Notice in particular the assembler directives '.data' and '.text' as well as the declarations of program constants. Understanding the basic flow of the example program will help you to complete your SPIM assignment quickly and painlessly. Once you have created the file and reviewed the code, it is time to move onto the next section.

13.4　　　QtSpim – Loading and Executing Programs

After the QtSpim application installation has been complete and the sample program has been created, you can execute the program to view the results. The use of QtSpim is described in the following sections.

13.4.1　　　Starting QtSpim

For Windows, this is typically performed with the standard "Start Menu -> Programs -> QtSpim" operation. For macOS, enter LaunchPad and click on QtSPim. For Linux, find the QtSpim icon (location is OS distribution dependent) and click on QtSpim.

13.4.2   Main Screen

The initial QtSpim screen will appear as shown below. There will be some minor differences based on the specific Operating System being used.

13.4.3   Load Program

To load the example program (and all programs), you can select the standard "**File → Reinitialize and Load File**" option from the menu bar. However, it is typically easier to select the **Reinitialize and Load File Icon** from the main screen (second file icon on the top left side).

Appendix B – QtSpim Tutorial

Note, the Load File option can be used on the initial load, but subsequent file loads will need to use the Reinitialize and Load File to ensure the appropriate reinitialization occurs.

Reinitialize and Load File Icon

Once selected, a standard open file dialog box will be displayed. Find and select 'asst0.asm' file (or whatever you named it) created in section 3.0.

Navigate as appropriate to find the example file previously created. When found, select the file (it will be highlighted) and click Open button (lower right hand corner).

The assembly process occurs as the file is being loaded. As such, any assembly syntax errors (i.e., misspelled instructions, undefined variables, etc.) are caught at this point. An appropriate error message is provided with a reference to the line number that caused the error.

When the file load is completed with no errors, the program is ready to run, but has not yet been executed. The screen will appear something like the following image.
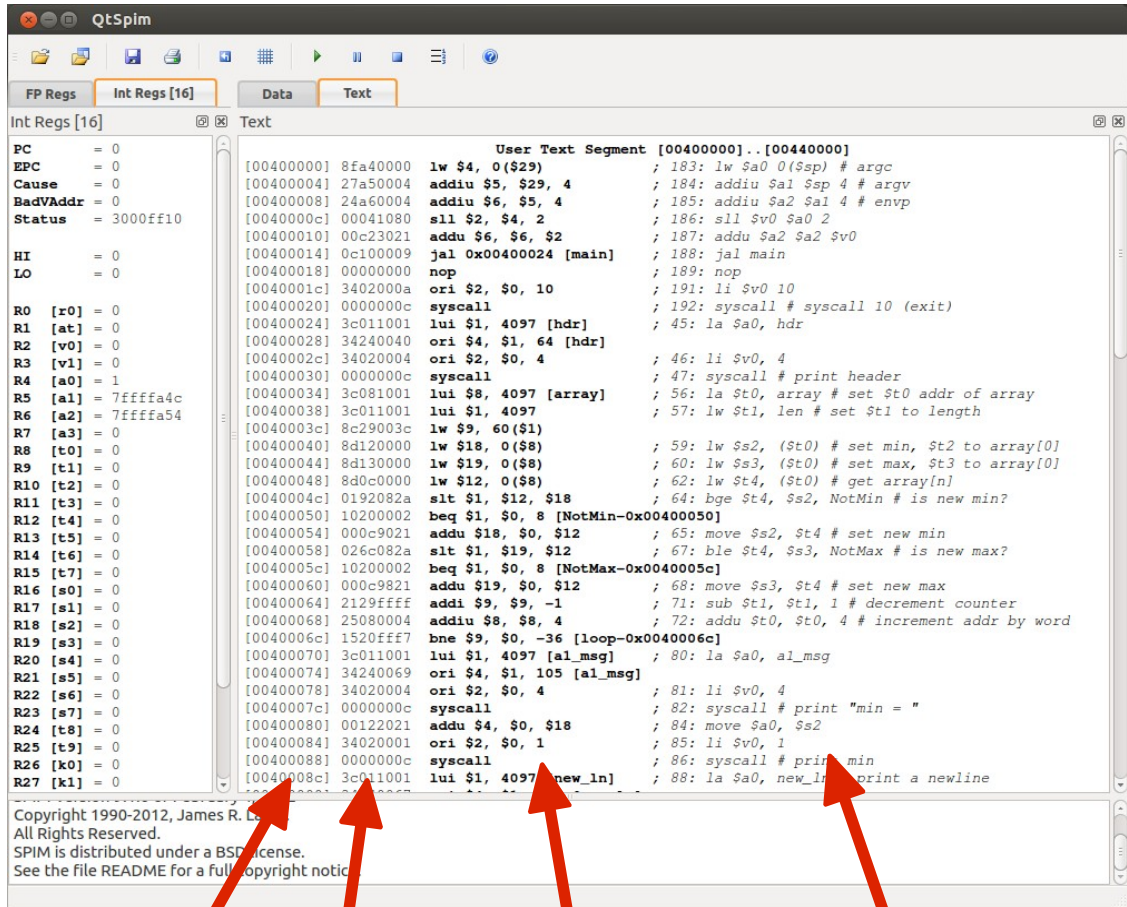
Addresses   OpCodes   Bare-Instructions   Pseudo-Instructions

The code is placed in Text Window. The first column of hex values (in the []'s) is the address of that line of code. The next hex value is the OpCode or hex value of the 1's and 0's that the CPU understands to be that instruction.

MIPS includes pseudo-instructions. That is an instruction that the CPU does not execute, but the programmer is allowed to use. The assembler, QtSpim here, accepts the instruction and inserts the real or *bare* instruction as appropriate.

13.4.4   Data Window

The data segment contains the data declared by your program (if any). To view the data segment, click on the Data Icon. The data window will appear similar to the following:

| Addresses | Data (Hex Representation) | Data (ASCII Representation) |

As before, the addresses are shown on the left side (with the []'s). The values at that address are shown in hex (middle) and in ASCII (right side). Depending on the specific type of data declarations, it may be easier to view the hex representation (i.e., like the

array of numbers from the example code) or the ASCII representation (i.e., the declared strings).

Note, right clicking in the Data Window will display a menu allowing the user to change the default hex representation to decimal representation (if desired).

13.4.5   Program Execution

To execute the entire program (uninterrupted), you can select the standard "**Simulator → Run/Continue**" option from the menu bar. However, it is typically easier to select the **Run/Continue Icon** from the main screen or to type the **F5** key.

Once typed, the program will be executed.

If a program performs input and/or output, it will be directed to the Console window.

Appendix B – QtSpim Tutorial

For example, the sample program (from Appendix A) will display the following in the Console window when executed.

For the sample program and the initial data set, these are the correct results.

13.4.6   Log File

QtSpim can create a log file saving and documenting the program results. To create a log file, you can select the standard "**File → Save Log File**" option from the menu bar. However, it is typically easier to select the **Save Log File Icon** from the main screen.
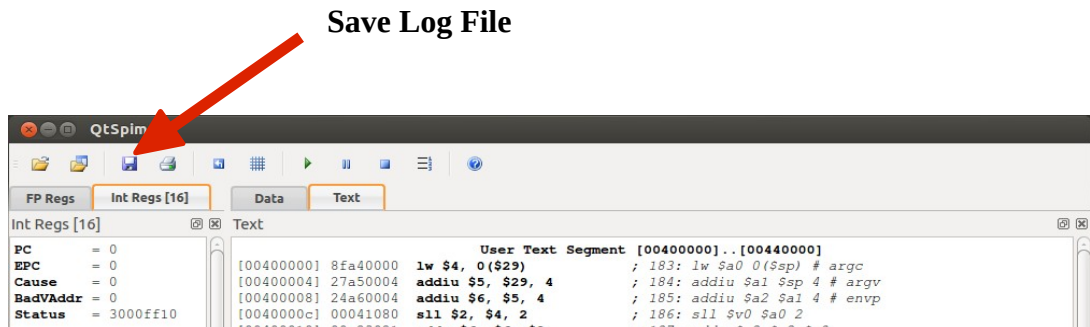
When selected, the Save Windows to Log File dialog box will be displayed as shown below on the left.

In general, the Text Segments and Console options should be selected as shown on the left.

Additionally, there is no default file name or location (for the log file). As such, a file name must be entered before it can be saved. This can be done by manually entering the name in the Save to file box or by selecting the **...** box (on the lower right side).

Appendix B – QtSpim Tutorial

When the **…** option is selected, a Save to Log File dialog box is displayed allowing selection of a location and the entry of a file name.
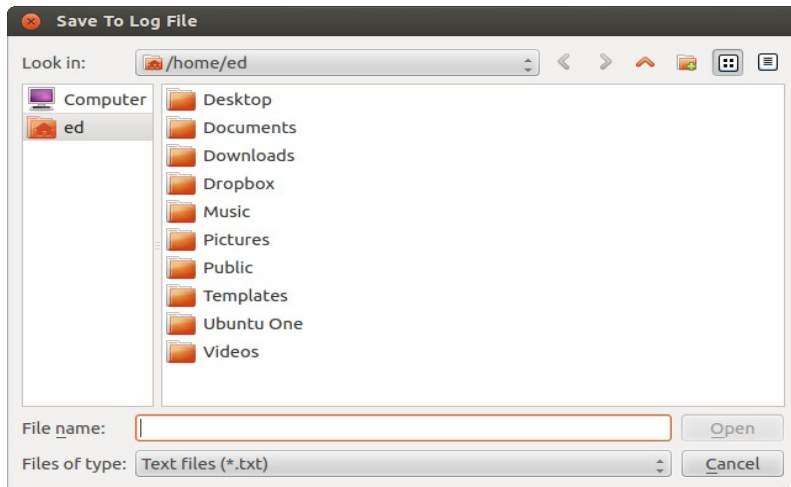
When completed correctly, the Save Windows To Log File box will appear similar the below image.
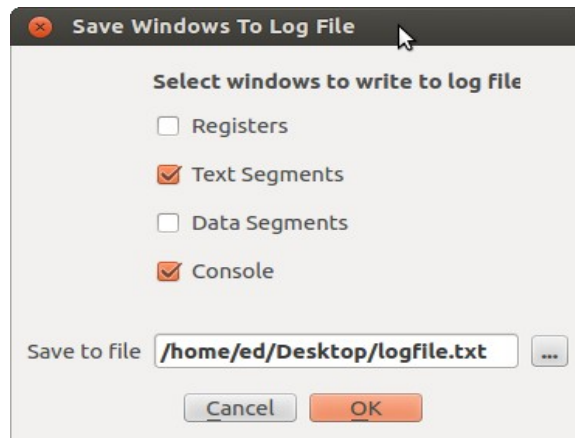
When the options are selected and the file name entered, the OK box can be selected which will save the log file.

13.4.7   Making Updates

In the highly unlikely event that the program does not work the first time or the program requirements are changed, the source file will need to be updated in a text editor. After the program source file is updated, it must be explicitly reloaded into QtSpim. The Reinitialize and Load File option must be used as described in section 13.4.3. Every change made to the source file must be re-loaded into QtSpim.

Once re-loaded, the program can be re-executed as noted in section 13.4.5. Refer to section 5.0 for information regarding debugging and controlling program execution.

13.5   Debugging

Often, looking at program source code will not help to find errors. The first step in debugging is to ensure that the file assembles correctly (or "reads" in the specific case of QtSpim). However, even if the file assembles, it still may not work correctly. In this case, the program must be debugged. In a broad sense, debugging is comparing the expected program results to actual program results. This requires a solid understanding of what the program is supposed to do and the specific order in which it does it → that is understanding the algorithm being used to solve the program. The algorithm should be noted in the program comments and can be used as a checklist for the debugging process.

Appendix B – QtSpim Tutorial

One potentially useful way to check the program status is to view the register contents. The current register contents are shown in registers window (left side) as shown in the image below.

Register Window
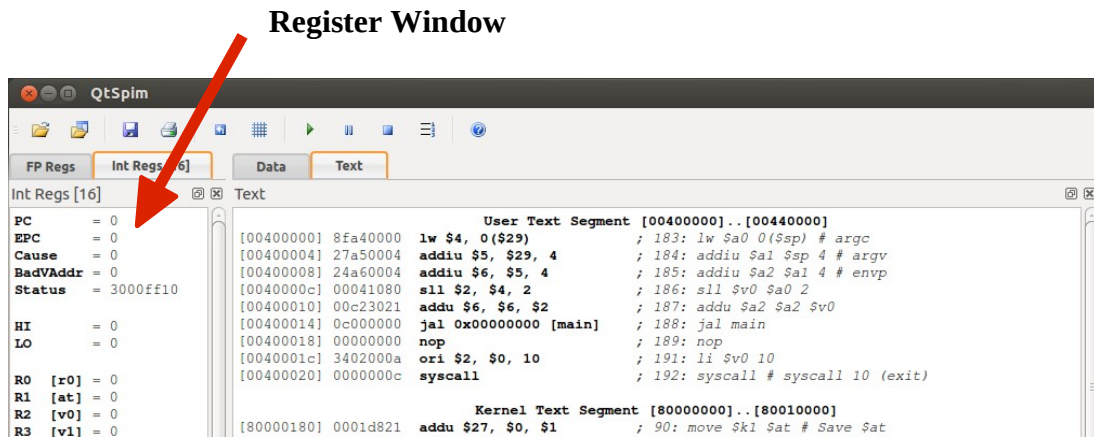
The overall debugging process can be simplified by using the QtSpim controlled execution functions. These functions include single stepping through the program and using one or more breakpoints. A breakpoint is a programmer selected location in the program where execution will be paused. When the program is paused the current program status can be checked by viewing the register contents and/or the data segment. Typically, a breakpoint will be set, the program executed (to that point), and from there single stepping through the program watching execution and checking the results (via register contents and/or data segment).

When stepping through the program, the *next instruction to be executed* is highlighted. As such, that instruction has **not** yet been executed. This highlighting is how to track the progress of the program execution.

To set a breakpoint, select an appropriate location. This should be chosen with a specific expectation in mind. For example, if a program does not produce the correct average for a list of numbers, a typical debugging strategy would be to see if the sum is correct (as it is required for the average calculation). As such, a breakpoint could be set after the loop and before the average calculation.

As an example, to set a breakpoint after the loop in the sample program (from Appendix A), the first instruction after the loop can be found in the Text Window. This will require looking at the pseudo-instructions (on the right side of the Text Window).

The first instruction after the loop in the example program is highlighted in orange (for reference) in the image below.

Note, the orange highlighting was added to this document for reference and will not be displayed in QtSpim during normal execution.

Appendix B – QtSpim Tutorial

When an appropriate instruction is determined, move the cursor to the instruction
address and right-click. The right-click will display the breakpoint menu as shown in
the image below.
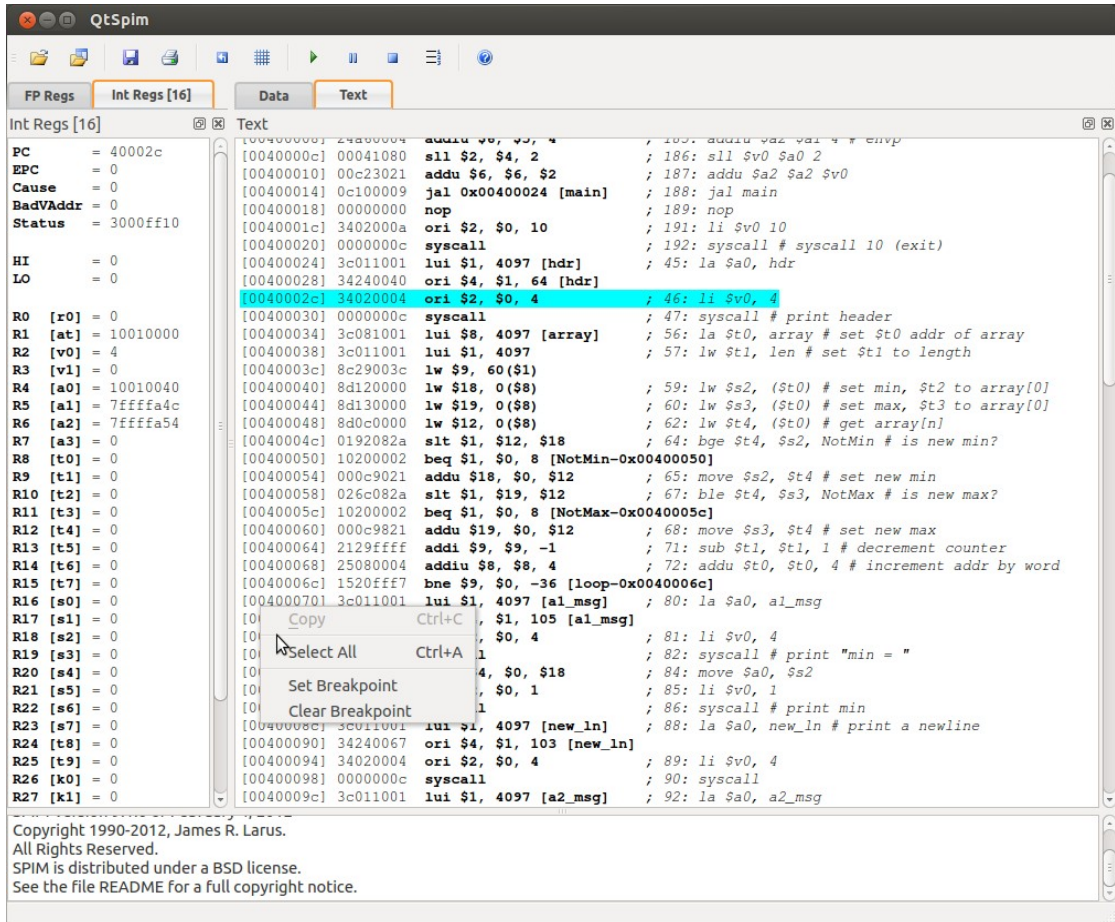
To set a breakpoint, select the Set Breakpoint option. If a breakpoint has already been
set, it can be cleared by selecting the Clear Breakpoint option.

Once the breakpoint has been set, it will be highlighted with a small red icon such as an **N** as shown in the following image. *Note*, different operating systems may use a different icon.

Select the Run/Continue option (as described in section 13.4.5) which will execute the program up to the selected breakpoint.

Appendix B – QtSpim Tutorial

When program execution reaches the breakpoint, it will be paused and a Breakpoint
dialog box displayed as shown in the below image.

The program execution can be halted by selecting the Abort box. The breakpoint can be
ignored, thus continuing to the next breakpoint or program termination, whichever
comes first.

However, typically the Single Step box will be selected upon entering the single step mode. The following image shows the result of selecting Single Step. Note, the highlighted instruction represents the next instruction to be executed and thus has not yet been executed.

Appendix B – QtSpim Tutorial

14.0   Appendix C – MIPS Instruction Set

This appendix presents a summary of the MIPS instructions as implemented within the QtSpim simulator. The instructions are grouped by like-operations and presented alphabetically.

The following table summarizes the notational conventions used.

Operand Notation	Description
Rdest	Destination operand. Must be a register. Since it is a destination operand, the contents will be over written with the new result.
FRdest	Destination operand. Must be a floating-point register. Since it is a destination operand, the contents will be over written with the new result.
Rsrc	Source operand. Must be a register. Register value is unchanged.
FRsrc	Source operand. Must be a floating-point register. Register value is unchanged.
Src	Source operand. Must be a register or an immediate value. Value is unchanged.
Imm	Immediate value
Mem	Memory location. May be a variable name or an indirect reference.

Refer to the chapter on Addressing Modes for more information regarding indirection.

14.1   Arithmetic Instructions

Below are a summary of the basic integer arithmetic instructions.

abs   Rdest, Rsrc
Absolute Value
Sets Rdest = absolute value of integer in Rsrc

add   Rdest, Rsrc, Src
Addition (with overflow)
Sets Rdest = Rsrc + Src (or imm)

addu   Rdest, Rsrc, Src
Addition (without overflow)
Sets Rdest = Rsrc + Src (or imm)

div   Rsrc1, Rsrc2
Divide (with overflow)
Set **$lo** = Rsrc / Src (or imm)
Remainder is placed in $hi

divu   Rsrc1, Rsrc2
Divide (without overflow)
Set **$lo** = Rsrc / Src (or imm)
Remainder is placed in $hi

div   Rdest, Rsrc, Src
Divide (with overflow)
Sets: Rdest = Rsrc / Src (or imm)

divu   Rdest, Rsrc, Src
Divide (without overflow)
Sets: Rdest = Rsrc / Src (or imm)

mul   Rdest, Rsrc, Src
Multiply (without overflow)
Sets: Rdest = Rsrc (Src (or imm)

mulo   Rdest, Rsrc, Src
Multiply (with overflow)
Sets: Rdest = Rsrc * Src (or imm)

mulou Rdest, Rsrc, Src	Unsigned Multiply (with overflow) Sets: **$lo** = Rsrc * Src (or imm)	
mult Rsrc1, Rsrc2	Multiply Sets **$hi**:**$lo** = Rsrc / Src (or imm)	
multu Rsrc1, Rsrc2	Unsigned Multiply Sets **$hi**:**$lo** = Rsrc / Src (or imm)	
neg Rdest, Rsrc	Negate Value (with overflow) Rdest = negative of integer in register Rsrc	
rem Rdest, Rsrc, Src	Remainder after division Rdest = remainder from Rsrc / Src (or imm)	
remu Rdest, Rsrc, Src	Unsigned Remainder Rdest = remainder from Rsrc / Src (or imm)	
sub Rdest, Rsrc, Src	Subtract (with overflow) Rdest = Rsrc – Src (or imm)	
subu Rdest, Rsrc, Src	Subtract (without overflow) Rdest = Rsrc – Src (or imm)	

14.2   Comparison Instructions

Below is a summary of the compare and set instructions. Programmers generally use the conditional branch and jump instructions as detailed in the next section.

seq   Rdest, Rsrc1, Src2   Set Equal
- Sets register Rdest to 1 if register Rsrc1 equals Src2 and to 0 otherwise

sge   Rdest, Rsrc1, Src2   Set Greater Than Equal
- Sets register Rdest to 1 if register Rsrc1 is greater than or equal Src2 and to 0 otherwise

sgeu   Rdest, Rsrc1, Src2   Set Greater Than Equal, Unsigned
- Sets register Rdest to 1 if register Rsrc1 is greater than or equal to Src2 and to 0 otherwise

sgt   Rdest, Rsrc1, Src2   Set Greater Than
- Sets register Rdest to 1 if register Rsrc1 is greater than Src2 and to 0 otherwise

sgtu   Rdest, Rsrc1, Src2   Set Greater Than, Unsigned
- Sets register Rdest to 1 if register Rsrc1 is greater than Src2 and to 0 otherwise

sle   Rdest, Rsrc1, Src2   Set Less Than Equal
- Sets register Rdest to 1 if register Rsrc1 is less than or equal to Src2 and to 0 otherwise

sleu   Rdest, Rsrc1, Src2   Set Less Than Equal, Unsigned
- Sets register Rdest to 1 if register Rsrc1 is less than or equal to Src2 and to 0 otherwise

slt   Rdest, Rsrc1, Src2   Set Less Than
- Sets register Rdest to 1 if register Rsrc1 is less than to Src2 and to 0 otherwise

slti Rdest, Rsrc1, Imm   Set Less Than, Immediate
- Sets register Rdest to 1 if register Rsrc1 is less than or equal to Imm and to 0 otherwise

sltu Rdest, Rsrc1, Src2   Set Less Than, Unsigned
- Sets register Rdest to 1 if register Rsrc1 is less than to Src2 and to 0 otherwise

sltiu Rdest, Rsrc1, Imm   Set Less Than Unsigned, Immediate
- Sets register Rdest to 1 if register Rsrc1 is less than Src2 (or Imm) and to 0 otherwise

sne   Rdest, Rsrc1, Src2   Set Not Equal
- Sets register Rdest to 1 if register Rsrc1 is not equal to Src2 and to 0 otherwise

14.3   Branch and Jump Instructions

Below are a summary of the basic conditional branch and jump instructions.

b label   Branch instruction
- Unconditionally branch to the instruction at the label

bczt label   Branch Co-processor *z* True
- Conditionally branch to the instruction at the label if co-processor *z*'s condition flag is true (false)

bczf   label

Branch Co-processor *z* False
- Conditionally branch to the instruction at
the label if co-processor *z*'s condition flag
is true (false)

beq   Rsrc1, Src2, label

Branch on Equal
- Conditionally branch to the instruction at
the label if the contents of register Rsrc1
equals Src2

beqz   Rsrc, label

Branch on Equal Zero
- Conditionally branch to the instruction at
the label if the contents of Rsrc equals 0

bge   Rsrc1, Src2, label

Branch on Greater Than or Equal
- Conditionally branch to the instruction at
the label if the contents of register Rsrc1
are greater than or equal to Src2

bgeu   Rsrc1, Src2, label

Branch on G Than or Equal, Unsigned
- Conditionally branch to the instruction at
the label if the contents of register Rsrc1
are greater than or equal to Src2

bgez   Rsrc, label

Branch on Greater Than or Equal Zero
- Conditionally branch to the instruction at
the label if the contents of Rsrc are greater
than or equal to 0

bgezal Rsrc, label

Branch on Greater Than or Equal Zero
and Link
- Conditionally branch to the instruction at
the label if the contents of Rsrc are greater
than or equal to 0. Saves the address of
the next instruction in **$ra**

bgt   Rsrc1, Src2, label   Branch on Greater Than
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 is greater than Src2

bgtu   Rsrc1, Src2, label   Branch on Greater Than, Unsigned
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are greater than Src2

bgtz   Rsrc, label   Branch on Greater Than Zero
- Conditionally branch to the instruction at the label if the contents of Rsrc are greater than 0

ble   Rsrc1, Src2, label   Branch on Less Than or Equal
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are less than or equal to Src2

bleu   Rsrc1, Src2, label   Branch on Less Than or Equal, Unsigned
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are less than or equal to Src2

blez   Rsrc, label   Branch on Less Than or Equal Zero
- Conditionally branch to the instruction at the label if the contents of Rsrc are less than or equal to 0

blezal   Rsrc, label   Branch on Less Than Equal or Zero And Link
- Conditionally branch to the instruction at the label if the contents of Rsrc are greater or equal to 0 or less than 0, respectively. Saves the address of the next instruction in register **$ra**

bltzal Rsrc, label   Branch on Less Than And Link
- Conditionally branch to the instruction at the label if the contents of Rsrc are less than 0 or less than 0, respectively. Save the address of the next instruction in register $ra

blt   Rsrc1, Src2, label   Branch on Less Than
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are less than Src2

bltu   Rsrc1, Src2, label   Branch on Less Than, Unsigned
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are less than Src2

bltz   Rsrc, label   Branch on Less Than Zero
- Conditionally branch to the instruction at the label if the contents of Rsrc are less than 0

bne   Rsrc1, Src2, label   Branch on Not Equal
- Conditionally branch to the instruction at the label if the contents of register Rsrc1 are not equal to Src2

bnez   Rsrc, label   Branch on Not Equal Zero
- Conditionally branch to the instruction at the label if the contents of Rsrc are not equal to 0

j   label   Jump
- Unconditionally jump to the instruction at the label

jal	**label**	Jump and Link - Unconditionally jump to the instruction at the label or whose address is in register Rsrc. Saves the address of the next instruction in register **$ra**
jalr	**Rsrc**	Jump and Link Register - Unconditionally jump to the instruction at the label or whose address is in register Rsrc. Saves the address of the next instruction in register **$ra**
jr	**Rsrc**	Jump Register - Unconditionally jump to the instruction whose address is in register Rsrc

14.4   Load Instructions

Below are a summary of the basic load instructions.

la	**Rdest, address**	Load Address - Load computed *address*, not the contents of the location, into register Rdest
lb	**Rdest, address**	Load Byte - Load the byte at *address* into register Rdest. The byte is sign-extended by the **lb**, but not the *lbu*, instruction
lbu	**Rdest, address**	Load Unsigned Byte - Load the byte at *address* into register Rdest. The byte is sign-extended by the **lb**, but not the *lbu*, instruction
ld	**Rdest, address**	Load Double-Word - Load the 64-bit quantity at *address* into registers Rdest and Rdest + 1

lh   Rdest, address

Load Halfword
- Load the 16-bit quantity (halfword) at *address* into register Rdest. The halfword is sign-extended

lhu   Rdest, address

Load Unsigned Halfword
- Load the 16-bit quantity (halfword) at *address* into register Rdest. The halfword is not sign-extended

lw   Rdest, address

Load Word
- Load the 32-bit quantity (word) at *address* into register Rdest

lwcz Rdest, address

Load Word Co-processor *z*
- Load the word at address into register Rdest of co-processor z (0-3)

lwl   Rdest, address

Load Word Left
- Load the left bytes from the word at the possibly-unaligned *address* into register Rdest

lwr   Rdest, address

Load Word Right
- Load the right bytes from the word at the possibly-unaligned *address* into register Rdest

ulh   Rdest, address

Unaligned Load Halfword
- Load the 16-bit quantity (halfword) at the possibly-unaligned *address* into register Rdest. The halfword is sign-extended.

ulhu Rdest, address   Unaligned Load Halfword Unsigned
- Load the 16-bit quantity (halfword) at the possibly-unaligned *address* into register Rdest. The halfword is not sign-extended

ulw   Rdest, address   Unaligned Load Word
- Load the 32-bit quantity (word) at the possibly-unaligned *address* into register Rdest

li   Rdest, imm   Load Immediate
- Move the immediate imm into register Rdest

lui   Rdest, imm   Load Upper Immediate
- Load the lower halfword of the immediate *imm* into the upper halfword of register Rdest. The lower bits of the register is set to 0

14.5   Logical Instructions

Below are a summary of the basic logical instructions.

and   Rdest, Rsrc1, Src2   AND

andi   Rdest, Rsrc1, Imm   AND Immediate
- Put the logical AND of the integers from register Rsrc1 and Src2 (or Imm) into register Rdest

nor   Rdest, Rsrc1, Src2   NOR
- Put the logical NOR of the integers from register Rsrc1 and Src2 into register Rdest

Appendix C – MIPS Instruction Set

not   Rdest, Rsrc

NOT
- Put the bitwise logical negation of the integer from register Rsrc into register Rdest

or   Rdest, Rsrc1, Src2

OR
- Put the logical OR of the integers from register Rsrc1 and Src2 into register Rdest

ori   Rdest, Rsrc1, Imm

OR Immediate
- Put the logical OR of the integers from register Rsrc1 and Imm into register Rdest

rol   Rdest, Rsrc1, Src2

Rotate Left
- Rotate the contents of register Rsrc1 left by the distance indicated by Src2 and put the result in register Rdest

ror   Rdest, Rsrc1, Src2

Rotate Right
- Rotate the contents of register Rsrc1 left (right) by the distance indicated by Src2 and put the result in register Rdest

sll   Rdest, Rsrc1, Src2

Shift Left Logical
- Shift the contents of register Rsrc1 left by the distance indicated by Src2 and put the result in register Rdest

sra   Rdest, Rsrc1, Src2

Shift Right Arithmetic
- Shift the contents of register Rsrc1 right by the distance indicated by Src2 and put the result in register Rdest

srl   Rdest, Rsrc1, Src2

Shift Right Logical
- Shift the contents of register Rsrc1 right by the distance indicated by Src2 and put the result in register Rdest

xor   Rdest, Rsrc1, Src2   XOR
- Put the logical XOR of the integers from register Rsrc1 and Src2 into register Rdest

xori Rdest, Rsrc1, Imm   XOR Immediate
- Put the logical XOR of the integers from register Rsrc1 and Imm into register Rdest

14.6   Store Instructions

Below are a summary of the basic store instructions.

sb   Rsrc, address   Store Byte
- Store the low byte from register Rsrc at *address*

sd   Rsrc, address   Store Double-Word
- Store the 64-bit quantity in registers Rsrc and Rsrc + 1 at *address*

sh   Rsrc, address   Store Halfword
- Store the low halfword from register Rsrc at *address*

sw   Rsrc, address   Store Word
- Store the word from register Rsrc at *address*

swcz Rsrc, address   Store Word Co-processor *z*
- Store the word from register Rsrc of co-processor z at address

swl   Rsrc, address   Store Word Left
- Store the left bytes from register Rsrc at the possibly-unaligned *address*

swr Rsrc, address	Store Word Right - Store the right bytes from register Rsrc at the possibly-unaligned *address*
ush Rsrc, address	Unaligned Store Halfword - Store the low halfword from register Rsrc at the possibly-unaligned *address*
usw Rsrc, address	Unaligned Store Word - Store the word from register Rsrc at the possibly-unaligned *address*

14.7   Data Movement Instructions

Below are a summary of the basic data movement instructions. The data movement implies data movement between registers.

move Rdest, Rsrc	Move the contents of Rsrc to Rdest. - The multiply and divide unit produces its result in two additional registers, **$hi** and **$lo**. These instructions move values to and from these registers. The multiply, divide, and remainder instructions described above are pseudo-instructions that make it appear as if this unit operates on the general registers and detect error conditions such as divide by zero or overflow.
mfhi Rdest	Move from **$hi** - Move the contents of the hi register to register Rdest
mflo Rdest	Move from **$lo** - Move the contents of the lo register to register Rdest

mthi Rdest

Move to **$hi**
- Move the contents register Rdest to the hi register.
- *Note*, Co-processors have their own register sets. This instruction move values between these registers and the CPU's registers.

mtlo Rdest

Move to **$lo**
- Move the contents register Rdest to the lo register.
- *Note*, Co-processors have their own register sets. This instruction move values between these registers and the CPU's registers.

mfc1 Rdest, FRsrc

Move From Co-processor 1
- Move the contents of co-processor 1 float register FRsrc to CPU integer register Rdest

mfc1.d Rdest, FRsrc1

Move Double From Co-processor 1
- Move the contents of floating-point registers FRsrc1 and FRsrc1+1 to CPU integer registers Rdest and Rdest + 1

mtc1   Rsrc, FRdest

Move To Co-processor 1
- Move the contents of CPU integer register Rsrc to co-processor 1 float register FRdest

mtc1.d Rsrc, FRdest

Move To Co-processor 1
- Move the contents of CPU integer registers Rsrc and Rsrc+1 to co-processor 1 float registers Frdest and FRdest+1.

14.8   Floating-Point Instructions

The MIPS has a floating-point co-processor (numbered 1) that operates on single precision (32-bit) and double precision (64-bit) floating-point numbers. This co-processor has its own registers, which are numbered **$f0** - **$f31**. Because these registers are only 32-bits wide, two of them are required to hold doubles. To simplify matters, floating-point operations only use even-numbered registers - including instructions that operate on single floats. Values are moved in or out of these registers a word (32-bits) at a time by **lwc1**, **swc1**, **mtc1**, and **mfc1** instructions described above or by the **l.s**, **l.d**, **s.s**, and **s.d** pseudo-instructions described below. The flag set by floating-point comparison operations is read by the CPU with its **bc1t** and **bc1f** instructions. In all instructions below, FRdest, FRsrc1, FRsrc2, and FRsrc are floating-point registers.

abs.d FRdest, FRsrc	Floating-point Absolute Value, Double - Compute the absolute value of the floating-point double in register FRsrc and put it in register FRdest
abs.s FRdest, FRsrc	Floating-point Absolute Value, Single - Compute the absolute value of the floating-point single in register FRsrc and put it in register FRdest
add.d FRdest, FRsrc1, FRsrc2	Floating-point Addition, Double - Compute the sum of the floating-point doubles in registers FRsrc1 and FRsrc2 and put it in register FRdest
add.s FRdest, FRsrc1, FRsrc2	Floating-point Addition, Single - Compute the sum of the floating-point singles in registers FRsrc1 and FRsrc2 and put it in register FRdest
c.eq.d FRsrc1, FRsrc2	Compare Equal, Double - Compare the floating-point double in register FRsrc1 against the one in FRsrc2 and set the floating-point condition flag true if they are equal

c.eq.s FRsrc1, FRsrc2

Compare Equal, Single
- Compare the floating-point single in register FRsrc1 against the one in FRsrc2 and set the floating-point condition flag true if they are equal

c.le.d FRsrc1, FRsrc2

Compare Less Than or Equal, Double
- Compare the floating-point double in register FRsrc1 against the one in FRsrc2 and set the floating-point condition flag true if the first is less than or equal to the second

c.le.s FRsrc1, FRsrc2

Compare Less Than or Equal, Single
- Compare the floating-point single precision in register FRsrc1 against the one in FRsrc2 and set the floating-point condition flag true if the first is less than or equal to the second

c.lt.d FRsrc1, FRsrc2

Compare Less Than, Double
- Compare the floating-point double in register FRsrc1 against the one in FRsrc2 and set the condition flag true if the first is less than the second

c.lt.s FRsrc1, FRsrc2

Compare Less Than, Single
- Compare the floating-point single in register FRsrc1 against the one in FRsrc2 and set the condition flag true if the first is less than the second

cvt.d.s FRdest, FRsrc

Convert Single to Double
- Convert the single precision floating-point number in register FRsrc to a double precision number and put it in register FRdest

`cvt.d.w  FRdest, FRsrc`	Convert Integer to Double - Convert the integer in register FRsrc to a double precision number and put it in register FRdest
`cvt.s.d  FRdest, FRsrc`	Convert Double to Single - Convert the double precision floating-point number in register FRsrc to a single precision number and put it in register FRdest
`cvt.s.w  FRdest, FRsrc`	Convert Integer to Single - Convert the integer in register FRsrc to a single precision number and put it in register FRdest
`cvt.w.d  FRdest, FRsrc`	Convert Double to Integer - Convert the double precision floating-point number in register FRsrc to an integer and put it in register FRdest
`cvt.w.s  FRdest, FRsrc`	Convert Single to Integer - Convert the single precision floating-point number in register FRsrc to an integer and put it in register FRdest
`div.d    FRdest, FRsrc1,` `          FRsrc2`	Floating-point Divide, Double - Compute the quotient of the floating-point doubles in registers FRsrc1 and FRsrc2 and put it in register FRdest.
`div.s    FRdest, FRsrc1,` `          FRsrc2`	Floating-point Divide, Single - Compute the quotient of the floating-point singles in registers FRsrc1 and FRsrc2 and put it in register FRdest.

l.d   FRdest, address

Load Floating-point, Double
- Load the floating-point double at address into register FRdest

l.s   FRdest, address

Load Floating-point, Single
- Load the floating-point single at address into register FRdest

mov.d   FRdest, FRsrc

Move Floating-point, Double
- Move the floating-point double from register FRsrc to register FRdest

mov.s FRdest, FRsrc

Move Floating-point, Single
- Move the floating-point single from register FRsrc to register FRdest

mul.d FRdest, FRsrc1, FRsrc2

Floating-point Multiply, Double
- Compute the product of the floating-point doubles in registers FRsrc1 and FRsrc2 and put it in register FRdest

mul.s FRdest, FRsrc1, FRsrc2

Floating-point Multiply, Single
- Compute the product of the floating-point singles in registers FRsrc1 and FRsrc2 and put it in register FRdest

neg.d FRdest, FRsrc

Negate, Double
- Store the floating-point double in register FRdest at address

neg.s FRdest, FRsrc

Negate, Single
Store the floating-point single in register FRdest at address

s.d   FRdest, address

Store Floating-point Double
- Store the floating-point double in register FRdest at address

s.s FRdest, address	Store Floating-point, Single - Store the floating-point single in register FRdest at address
sub.d FRdest, FRsrc1, FRsrc2	Floating-point Subtract, Double - Compute the difference of the floating-point doubles in registers FRsrc1 and FRsrc2 and put it in register FRdest
sub.s FRdest, FRsrc1, FRsrc2	Floating-point Subtract, Single - Compute the difference of the floating-point singles in registers FRsrc1 and FRsrc2 and put it in register FRdest

14.9   Exception and Trap Handling Instructions

Below are a summary of the exception and trap instructions.

rfe	Return From Exception - Restore the Status register
syscall	System Call - Transfer control to system routine. Register **$v0** contains the number of the system call
break n	Break - Cause exception *n*. - *Note*, Exception 1 is reserved for the debugger
nop	No operation - Do nothing

15.0   Appendix D – ASCII Table

This appendix provides a copy of the ASCII Table for reference.

Char	Dec	Hex		Char	Dec	Hex		Char	Dec	Hex		Char	Dec	Hex
NUL	0	0x00		spc	32	0x20		@	64	0x40		`	96	0x60
SOH	1	0x01		!	33	0x21		A	65	0x41		a	97	0x61
STX	2	0x02		"	34	0x22		B	66	0x42		b	98	0x62
ETX	3	0x03		#	35	0x23		C	67	0x43		c	99	0x63
EOT	4	0x04		$	36	0x24		D	68	0x44		d	100	0x64
ENQ	5	0x05		%	37	0x25		E	69	0x45		e	101	0x65
ACK	6	0x06		&	38	0x26		F	70	0x46		f	102	0x66
BEL	7	0x07		'	39	0x27		G	71	0x47		g	103	0x67
BS	8	0x08		(	40	0x28		H	72	0x48		h	104	0x68
TAB	9	0x09		)	41	0x29		I	73	0x49		i	105	0x69
LF	10	0x0A		*	42	0x2A		J	74	0x4A		j	106	0x6A
VT	11	0x0B		+	43	0x2B		K	75	0x4B		k	107	0x6B
FF	12	0x0C		,	44	0x2C		L	76	0x4C		l	108	0x6C
CR	13	0x0D		-	45	0x2D		M	77	0x4D		m	109	0x6D
SO	14	0x0E		.	46	0x2E		N	78	0x4E		n	110	0x6E
SI	15	0x0F		/	47	0x2F		O	79	0x4F		o	111	0x6F
DLE	16	0x10		0	48	0x30		P	80	0x50		p	112	0x70
DC1	17	0x11		1	49	0x31		Q	81	0x51		q	113	0x71
DC2	18	0x12		2	50	0x32		R	82	0x52		r	114	0x72
DC3	19	0x13		3	51	0x33		S	83	0x53		s	115	0x73
DC4	20	0x14		4	52	0x34		T	84	0x54		t	116	0x74
NAK	21	0x15		5	53	0x35		U	85	0x55		u	117	0x75
SYN	22	0x16		6	54	0x36		V	86	0x56		v	118	0x76
ETB	23	0x17		7	55	0x37		W	87	0x57		w	119	0x77
CAN	24	0x18		8	56	0x38		X	88	0x58		x	120	0x78

Appendix D – ASCII Table

EM	25	0x19		9	57	0x39		Y	89	0x59		y	121	0x79
SUB	26	0x1A		:	58	0x3A		Z	90	0x5A		z	122	0x7A
ESC	27	0x1B		;	59	0x3B		[	91	0x5B		{	123	0x7B
FS	28	0x1C		<	60	0x3C		\	92	0x5C		|	124	0x7C
GS	29	0x1D		=	61	0x3D		]	93	0x5D		}	125	0x7D
RS	30	0x1E		>	62	0x3E		^	94	0x5E		~	126	0x7E
US	31	0x1F		?	63	0x3F		_	95	0x5F		DEL	127	0x7F

For additional information and a more complete listing of the ASCII codes (including the extended ASCII characters), refer to http://www.asciitable.com/

16.0   Alphabetical Index

Alphabetical Index